

ALBERT Sentiment Analysis Experiments on Sentiment

- ✓ 140 Dataset using baseline model (and with PEFT -- LoRA, BitFit, Prompt Tuning)

"""\nWe begin our process by installing packages such as pytorch, which is used extensively, transformers and datasets packages, which are used to run the ALBERT transformer !

```
!pip install torch transformers datasets -q
```

→

363.4/363.4 MB	3.2 MB/s	eta	0:00:00
13.8/13.8 MB	83.6 MB/s	eta	0:00:00
24.6/24.6 MB	53.4 MB/s	eta	0:00:00
883.7/883.7 kB	49.8 MB/s	eta	0:00:00
664.8/664.8 MB	1.6 MB/s	eta	0:00:00
211.5/211.5 MB	11.7 MB/s	eta	0:00:00
56.3/56.3 MB	43.2 MB/s	eta	0:00:00
127.9/127.9 MB	19.1 MB/s	eta	0:00:00
207.5/207.5 MB	4.8 MB/s	eta	0:00:00
21.1/21.1 MB	110.3 MB/s	eta	0:00:00
491.2/491.2 kB	43.3 MB/s	eta	0:00:00
116.3/116.3 kB	11.2 MB/s	eta	0:00:00
183.9/183.9 kB	17.8 MB/s	eta	0:00:00
143.5/143.5 kB	16.3 MB/s	eta	0:00:00
194.8/194.8 kB	19.3 MB/s	eta	0:00:00

ERROR: pip's dependency resolver does not currently take into account all the gcsfs 2025.3.2 requires fsspec==2025.3.2, but you have fsspec 2024.12.0 which

"""\nThis step configures the credentials of the active user to seamlessly enable p

```
!git config --global credential.helper store
```

""""We next import the installed packages, namely the ALBERT model""""

```
import torch
from torch.utils.data import DataLoader
from datasets import load_dataset
from transformers import AutoTokenizer, AutoModelForSequenceClassification, DataCollatorWithPadding

import time
from sklearn.metrics import classification_report
```

"""" We next instantiate (load) our temporary dataset, calling to our sentiment140.py

```
dataset = load_dataset("/content/sentiment140.py", name="sentiment140")
full_train = dataset["train"]

print("Train size:", len(dataset["train"]))
print("Test size:", len(dataset["test"]))
```

→ The repository for sentiment140 contains custom code which must be executed to. You can avoid this prompt in future by passing the argument `trust_remote_code`

Do you wish to run the custom code? [y/N] y

Downloading data: 100% 81.4M/81.4M [00:37<00:00, 5.42MB/s]

Generating train split: 1600000/0 [00:49<00:00, 32381.84 examples/s]

Generating test split: 498/0 [00:00<00:00, 18027.56 examples/s]

Train size: 1600000

Test size: 498

"""" We next import a few packages for randomization of our sampling, re for text

```
import random
import re
from datasets import Dataset
```

"""\ With the entire 1.6m entry dataset loaded in as `full_train` above, we next filter out the 1.6m entries with sentiment 3 (though there seemed not to be any such instances), and we define negative and positive samples. We lastly overwrite our dataset with just the 50k class-balanced records from previous steps. This is done for our resource-efficiency computationally-constrained focus.""""

```
all_data = [x for x in full_train if x["sentiment"] in [0,4]]  
  
negative = [x for x in all_data if x["sentiment"] == 0]  
positive = [x for x in all_data if x["sentiment"] == 4]  
  
random.seed(42)  
  
negative_sample = random.sample(negative, 25000)  
positive_sample = random.sample(positive, 25000)  
  
sampled_data = negative_sample + positive_sample  
  
random.shuffle(sampled_data)  
  
dataset = Dataset.from_list(sampled_data)
```

"""\ Next, with our 50k sentiment dataset, we perform pre-processing for standardizing the text. This includes removing noise in the form of mentions (e.g. @gatech), URLs (e.g. <https://...>), hashtags (e.g. #gatech), and non-alphanumeric characters (e.g. emojis, capitalization, punctuation).""""

```
def clean_text(text):  
    text = text.lower()  
    text = re.sub(r"http\S+", "", text)  
    text = re.sub(r"@[\w]+", "", text)  
    text = re.sub(r"#[\w]+", "", text)  
    text = re.sub(r"[\^a-z0-9\s]", "", text)  
    return text.strip()
```

```
dataset = dataset.map(lambda x: {"text": clean_text(x["text"])})
```



Map: 100%

50000/50000 [00:04<00:00, 14983.40 examples/s]

```
model_name = "albert-base-v2"
num_labels = 2
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForSequenceClassification.from_pretrained(model_name, num_labels)

→ /usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: Use
  The secret `HF_TOKEN` does not exist in your Colab secrets.
  To authenticate with the Hugging Face Hub, create a token in your settings tak
  You will be able to reuse this secret in all of your notebooks.
  Please note that authentication is recommended but still optional to access pu
    warnings.warn(
tokenizer_config.json: 100%                                         25.0/25.0 [00:00<00:00, 3.13kB/s]

config.json: 100%                                         684/684 [00:00<00:00, 73.8kB/s]

spiece.model: 100%                                         760k/760k [00:00<00:00, 3.50MB/s]

tokenizer.json: 100%                                         1.31M/1.31M [00:00<00:00, 5.92MB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, bu
model.safetensors: 100%                                         47.4M/47.4M [00:00<00:00, 194MB/s]

Some weights of AlbertForSequenceClassification were not initialized from the
You should probably TRAIN this model on a down-stream task to be able to use i
```

```
"""\n    We apply mapping from the seemingly arbitrary 0,4 scale to 0,1 for standard b\n\ndef map_labels(example):\n    example["sentiment"] = 0 if example["sentiment"] == 0 else 1\n    return example\n\ndataset = dataset.map(map_labels)\n\nprint("Dataset size:", len(dataset))\n\nfrom collections import Counter\nprint("Class distribution:", Counter(dataset["sentiment"]))\n\n→ Map: 100%                                         50000/50000 [00:03<00:00, 8267.03 examples/\n\n                                         s]\n\nDataset size: 50000
```

DOWNSAMPLING FOR TRAINING

"" We next perform downsampling and stratified splitting of the data 80/20 Train.

```
from sklearn.model_selection import train_test_split
from datasets import Dataset, DatasetDict
from collections import Counter
import pandas as pd

df = dataset.to_pandas()
df = df.dropna(subset=["text", "sentiment"])

print("Initial class distribution:", Counter(df["sentiment"]))

df_train, df_test = train_test_split(
    df,
    stratify=df["sentiment"],
    test_size=0.2,
    random_state=42
)

print("Train size:", len(df_train))
print("Test size:", len(df_test))
print("Train class distribution:", Counter(df_train["sentiment"]))
print("Test class distribution:", Counter(df_test["sentiment"]))

train_dataset = Dataset.from_pandas(df_train).remove_columns(["__index_level_0__"])
test_dataset = Dataset.from_pandas(df_test).remove_columns(["__index_level_0__"])

def tokenize(example):
    return tokenizer(example["text"], truncation=True, padding="max_length", max_)

train_dataset = train_dataset.map(tokenize, batched=True)
test_dataset = test_dataset.map(tokenize, batched=True)

train_dataset.set_format("torch", columns=["input_ids", "attention_mask", "sentiment"])
test_dataset.set_format("torch", columns=["input_ids", "attention_mask", "sentiment"])

tokenized_dataset = DatasetDict({
    "train": train_dataset,
    "test": test_dataset
})
```

```
Initial class distribution: Counter({1: 25000, 0: 25000})
Train size: 40000
Test size: 10000
Train class distribution: Counter({1: 20000, 0: 20000})
Test class distribution: Counter({0: 5000, 1: 5000})
Map: 100% 40000/40000 [00:03<00:00, 10037.48 examples/s]
```

Map: 100%

s]

10000/10000 [00:00>00:00, 10331.71 examples/s]

"" We print the head of each of the train/test sets to visualize our cleaned data.

```
print("\nSample training examples:")
display(df_train.head(5))
```

```
print("\nSample test examples:")
display(df_test.head(5))
```



Sample training examples:

		text	date	user	sentiment	query
47782	finally set up wireless internet huzzah for tw...		Sat May 30 23:43:31 PDT 2009	alexwilliamson	1	NO_QUERY
20407	lebron james please dont leave usfor the love ...		Sat May 30 20:58:04 PDT 2009	ryangett	0	NO_QUERY
42997	i broke our site		Sat May 30 23:21:22 PDT 2009	DjDATZ	0	NO_QUERY
19678	ugh idk if thats going to be possible my fri...		Fri May 22 08:28:11 PDT 2009	JulieAnnCook	0	NO_QUERY
13754	wrong place at the wrong time always sigh		Wed Jun 17 08:20:28 PDT 2009	Nadiahazman	0	NO_QUERY

Sample test examples:

		text	date	user	sentiment	query

```
""" We initialize our dataloader for each of the sets, fix their batch sizes  
and randomize their order"""
```

```
train_loader = DataLoader(tokenized_dataset["train"], batch_size=16, shuffle=True  
test_loader = DataLoader(tokenized_dataset["test"], batch_size=16)
```

```
""" We define our optimizer using Adam and set a conservative learning rate and  
weight decay (though later hyperparameter search will overwrite)"""
```

```
optimizer = torch.optim.AdamW(  
    filter(lambda p: p.requires_grad, model.parameters()),  
    lr=5e-5,  
    weight_decay=0.01  
)  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
model.to(device)
```

```

→ AlbertForSequenceClassification(
    (albert): AlbertModel(
        (embeddings): AlbertEmbeddings(
            (word_embeddings): Embedding(30000, 128, padding_idx=0)
            (position_embeddings): Embedding(512, 128)
            (token_type_embeddings): Embedding(2, 128)
            (LayerNorm): LayerNorm((128,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0, inplace=False)
        )
        (encoder): AlbertTransformer(
            (embedding_hidden_mapping_in): Linear(in_features=128,
            out_features=768, bias=True)
            (albert_layer_groups): ModuleList(
                (0): AlbertLayerGroup(
                    (albert_layers): ModuleList(
                        (0): AlbertLayer(
                            (full_layer_layer_norm): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
                            (attention): AlbertSdpAttention(
                                (query): Linear(in_features=768, out_features=768, bias=True)
                                (key): Linear(in_features=768, out_features=768, bias=True)
                                (value): Linear(in_features=768, out_features=768, bias=True)
                                (attention_dropout): Dropout(p=0, inplace=False)
                                (output_dropout): Dropout(p=0, inplace=False)
                                (dense): Linear(in_features=768, out_features=768, bias=True)
                                (LayerNorm): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
                            )
                            (ffn): Linear(in_features=768, out_features=3072, bias=True)
                            (ffn_output): Linear(in_features=3072, out_features=768,
bias=True)
                            (activation): NewGELUActivation()
                            (dropout): Dropout(p=0, inplace=False)
                        )
                    )
                )
            )
        )
        (pooler): Linear(in_features=768, out_features=768, bias=True)
        (pooler_activation): Tanh()
    )
    (dropout): Dropout(p=0.1, inplace=False)
    (classifier): Linear(in_features=768, out_features=2, bias=True)
)

```

""" Baseline inference for binary sentiment analysis task run on ALBERT without PEFT (i.e. without BitFit and/or LoRA)"""

```
import time
import torch
from sklearn.metrics import classification_report, confusion_matrix, f1_score
import seaborn as sns
import matplotlib.pyplot as plt
from torch import autocast

inference_start = time.time()

model.eval()
total_correct = 0
total_samples = 0
all_preds = []
all_labels = []

with torch.no_grad():
    for batch in test_loader:
        input_ids = batch["input_ids"].to(device)
        attention_mask = batch["attention_mask"].to(device)
        labels = batch["sentiment"].to(device)

        with autocast(device_type='cuda'):
            outputs = model(input_ids=input_ids, attention_mask=attention_mask)
            logits = outputs.logits
            predictions = torch.argmax(logits, dim=-1)

        all_preds.extend(predictions.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

        total_correct += (predictions == labels).sum().item()
        total_samples += labels.size(0)

accuracy = total_correct / total_samples
f1_macro = f1_score(all_labels, all_preds, average="macro")
f1_weighted = f1_score(all_labels, all_preds, average="weighted")
inference_time = time.time() - inference_start

print(f'\nBaseline Inference Performance - ALBERT on Sentiment140\n')
print(f"\nTest Accuracy : {accuracy:.4f}")
print(f"F1 Score (macro): {f1_macro:.4f}")
print(f"F1 Score (weighted): {f1_weighted:.4f}")
print(f"Inference Time : {inference_time:.2f}s")
print("\nClassification Report:")
print(classification_report(all_labels, all_preds, target_names=["Negative", "Pos"])
```

```
cm = confusion_matrix(all_labels, all_preds)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=["Negative", "Positive"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix – Baseline Inference (ALBERT on Sentiment140)")
plt.show()
```



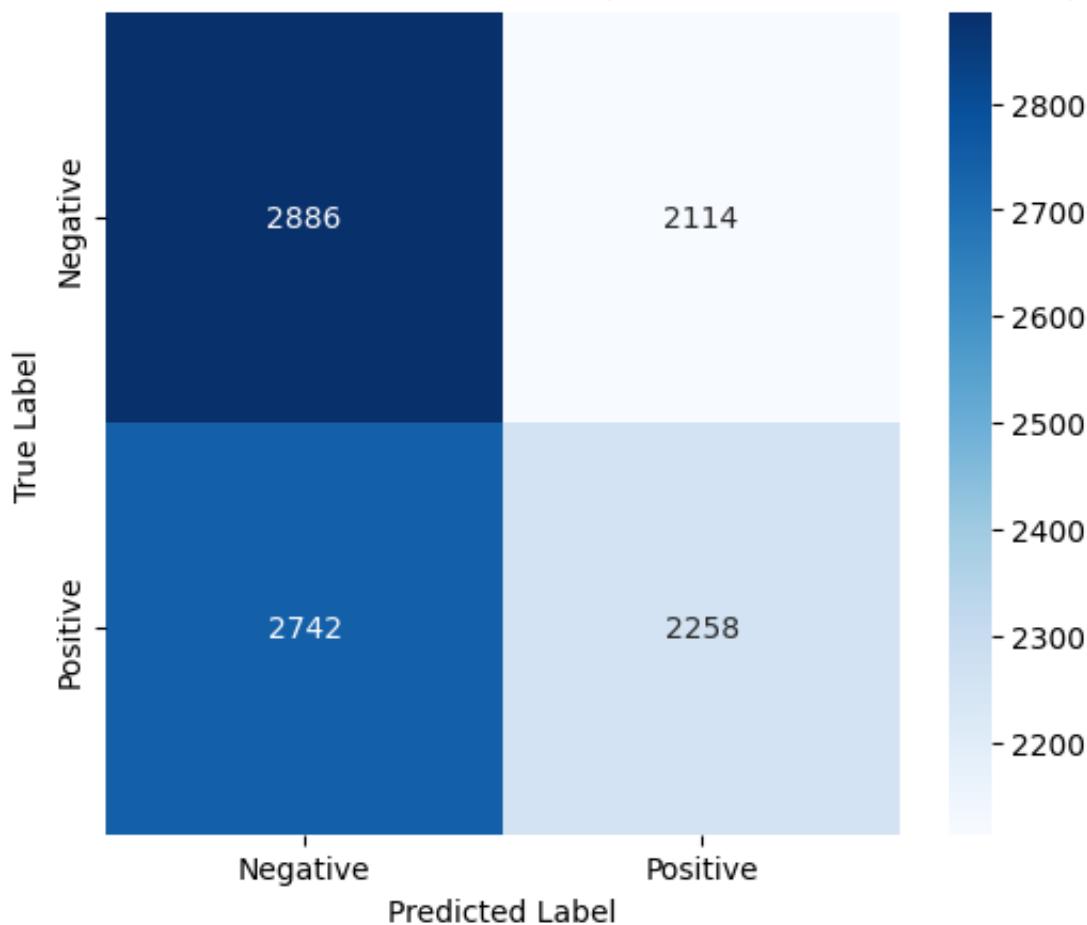
Baseline Inference Performance - ALBERT on Sentiment140

```
Test Accuracy : 0.5144
F1 Score (macro): 0.5125
F1 Score (weighted): 0.5125
Inference Time : 19.84s
```

Classification Report:

	precision	recall	f1-score	support
Negative	0.51	0.58	0.54	5000
Positive	0.52	0.45	0.48	5000
accuracy			0.51	10000
macro avg	0.51	0.51	0.51	10000
weighted avg	0.51	0.51	0.51	10000

Confusion Matrix - Baseline Inference (ALBERT on Sentiment140)



✓ LORA

```
""" Install Parameter Efficient Finetuning Packages (e.g. LoRA and BitFit)"""

!pip install peft -q

""" Importing LoRA packages """

import gc
import torch
import time
import pandas as pd
from tqdm import tqdm
from transformers import AutoModelForSequenceClassification, AutoTokenizer, DataC
from peft import get_peft_model, LoraConfig, TaskType
from sklearn.metrics import classification_report, f1_score
from torch.utils.data import DataLoader

""" LoRA parameter setup """

learning_rates = [5e-5, 1e-4]
batch_sizes = [8, 16]
epochs = 6

""" Training on ALBERT model using LoRA and output dataset generation (saved as ..

results = []

for lr in learning_rates:
    for batch_size in batch_sizes:
        print(f"Running LoRA with LR={lr}, batch_size={batch_size}")

        # loading ALBERT model
        model_name = "albert-base-v2"
        tokenizer = AutoTokenizer.from_pretrained(model_name)
        model = AutoModelForSequenceClassification.from_pretrained(model_name, nu

        # LoRA param update config
        lora_config = LoraConfig(
            task_type=TaskType.SEQ_CLS,
```

```
r=16,  
lora_alpha=32,  
lora_dropout=0.1,  
bias="none",  
target_modules=["query", "key", "value"]  
)  
  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
model.to(device)  
  
# instantiate dataloader  
data_collator = DataCollatorWithPadding(tokenizer)  
train_dataloader = DataLoader(tokenized_dataset["train"], batch_size=batch_size)  
test_dataloader = DataLoader(tokenized_dataset["test"], batch_size=batch_size)  
  
# adam optimizer  
optimizer = torch.optim.AdamW(filter(lambda p: p.requires_grad, model.parameters()))  
  
# begin training  
model.train()  
start_time = time.time()  
epoch_logs = []  
  
for epoch in range(1, epochs + 1):  
    running_loss = 0.0  
    correct = 0  
    total = 0  
    loop = tqdm(train_dataloader, leave=False)  
    for step, batch in enumerate(loop):  
        batch = {  
            "input_ids": batch["input_ids"].to(device),  
            "attention_mask": batch["attention_mask"].to(device),  
            "labels": batch["sentiment"].to(device)  
        }  
        with autocast(device_type='cuda'):  
            outputs = model(**batch)  
            loss = outputs.loss  
            preds = torch.argmax(outputs.logits, dim=1)  
            correct += (preds == batch['labels']).sum().item()  
            total += batch['labels'].size(0)  
  
            optimizer.zero_grad()  
            loss.backward()  
            optimizer.step()
```

```
        running_loss += loss.item()

    avg_train_loss = running_loss / (step + 1)
    train_accuracy = correct / total

    # perform per epoch evaluation
    model.eval()
    val_running_loss = 0.0
    y_true, y_pred = [], []
    inference_start = time.time()
    with torch.no_grad():
        with autocast(device_type='cuda'):
            for batch in test_dataloader:
                batch = {
                    "input_ids": batch["input_ids"].to(device),
                    "attention_mask": batch["attention_mask"].to(device),
                    "labels": batch["sentiment"].to(device)
                }
                outputs = model(**batch)
                preds = torch.argmax(outputs.logits, dim=1)
                y_true.extend(batch["labels"].cpu().numpy())
                y_pred.extend(preds.cpu().numpy())
                val_running_loss += outputs.loss.item()

    avg_val_loss = val_running_loss / len(test_dataloader)
    inference_time = time.time() - inference_start

    report = classification_report(y_true, y_pred, output_dict=True)
    val_accuracy = report["accuracy"]
    val_f1 = report["weighted avg"]["f1-score"]

    epoch_logs.append({
        "epoch": epoch,
        "lr": lr,
        "batch_size": batch_size,
        "train_loss": avg_train_loss,
        "train_accuracy": train_accuracy,
        "val_loss": avg_val_loss,
        "val_accuracy": val_accuracy
    })

    if epoch == epochs:
        total_correct = sum(yt == yp for yt, yp in zip(y_true, y_pred))
        total_samples = len(y_true)
        accuracy = total_correct / total_samples
```

```
f1_macro = f1_score(y_true, y_pred, average="macro")
f1_weighted = f1_score(y_true, y_pred, average="weighted")

print(f"\n[Final Epoch {epoch}] Inference Metrics:")
print(f"Test Accuracy      : {accuracy:.4f}")
print(f"F1 Score (macro)  : {f1_macro:.4f}")
print(f"F1 Score (weighted): {f1_weighted:.4f}")
print(f"Inference Time     : {inference_time:.2f} seconds")
print("\nClassification Report: ALBERT w/ LoRA on Sentiment140\n")
print(classification_report(y_true, y_pred, target_names=["Negative", "Positive"]))

model.train()

end_time = time.time()
training_time = end_time - start_time

# begin datalogging per lr/bs
epoch_logs_df = pd.DataFrame(epoch_logs)
epoch_logs_df.to_csv(f"sent_albert_lora_epoch_logs_lr{lr}_bs{batch_size}.csv")

# save inference metrics per lr/bs
metrics_summary_df = pd.DataFrame(report).transpose()
metrics_summary_df.to_csv(f"sent_albert_lora_inference_metrics_summary_lr{lr}_bs{batch_size}.csv")

# save inference predictions for the final epoch
predictions_df = pd.DataFrame({
    "y_true": y_true,
    "y_pred": y_pred
})
predictions_df.to_csv(f"sent_albert_lora_inference_predictions_lr{lr}_bs{batch_size}.csv")

# log memory usage
max_memory = torch.cuda.max_memory_allocated() / (1024 ** 3) if torch.cuda.is_available() else None
print(f"Max Memory Usage: {max_memory} GB")

# save model params and metrics
results.append({
    "method": "LoRA",
    "learning_rate": lr,
    "batch_size": batch_size,
    "accuracy": val_accuracy,
    "f1": val_f1,
    "training_time": training_time,
    "inference_time": inference_time,
    "max_memory": max_memory
}))
```

```
# empty cache to conserve compute
del model, tokenizer, optimizer
torch.cuda.empty_cache()
gc.collect()

# ranked performance by val acc
results = sorted(results, key=lambda x: x["accuracy"], reverse=True)

# save overall results
results_df = pd.DataFrame(results)
results_df.to_csv("sent_albert_lora_results.csv", index=False)

# save best final config and metrics
final_summary_df = pd.DataFrame({
    "Method": ["LoRA"],
    "Best LR": [results[0]["learning_rate"]],
    "Best Batch Size": [results[0]["batch_size"]],
    "Accuracy": [results[0]["accuracy"]],
    "F1 Score": [results[0]["f1"]],
    "Training Time (s)": [results[0]["training_time"]],
    "Inference Time (s)": [results[0]["inference_time"]],
    "Max GPU Memory (GB)": [results[0]["max_memory"]]
})
final_summary_df.to_csv("sent_albert_lora_final_comparison_lora.csv", index=False)

print("All LoRA Grid Search Results:")
for r in results:
    print(r)

print("\nBest LoRA Configuration:")
print(results[0])

→ /usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:150
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

[Final Epoch 6] Inference Metrics:
Test Accuracy      : 0.5000
F1 Score (macro)  : 0.3333
F1 Score (weighted): 0.3333
Inference Time     : 17.62 seconds
```

Classification Report: ALBERT w/ LoRA on Sentiment140

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	5000
Positive	0.50	1.00	0.67	5000
accuracy			0.50	10000
macro avg	0.25	0.50	0.33	10000
weighted avg	0.25	0.50	0.33	10000

Running LoRA with LR=0.0001, batch_size=16

```
lora_best_lr = results[0]["learning_rate"]
lora_best_bs = results[0]["batch_size"]
```

```
# Construct filename
best_report_file = f"sent_albert_lora_inference_metrics_summary_lr{lora_best_lr}_"

# Load the saved best report
best_report_df = pd.read_csv(best_report_file)
print("\nClassification Report for Best Configuration:")
print(best_report_df)

best_preds_df = pd.read_csv(f"sent_albert_lora_inference_predictions_lr{lora_best_lr}_")
print("\nInference Predictions for Best Configuration:")
print(best_preds_df)

y_true = best_preds_df["y_true"]
y_pred = best_preds_df["y_pred"]

cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=["Negative", "Positive"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - ALBERT w/ LoRA on Sentiment140 \n")
plt.show()
```



Classification Report for Best Configuration:

	Unnamed: 0	precision	recall	f1-score	support
0	0	0.739978	0.8196	0.777757	5000.0000
1	1	0.797848	0.7120	0.752484	5000.0000
2	accuracy	0.765800	0.7658	0.765800	0.7658
3	macro avg	0.768913	0.7658	0.765120	10000.0000
4	weighted avg	0.768913	0.7658	0.765120	10000.0000

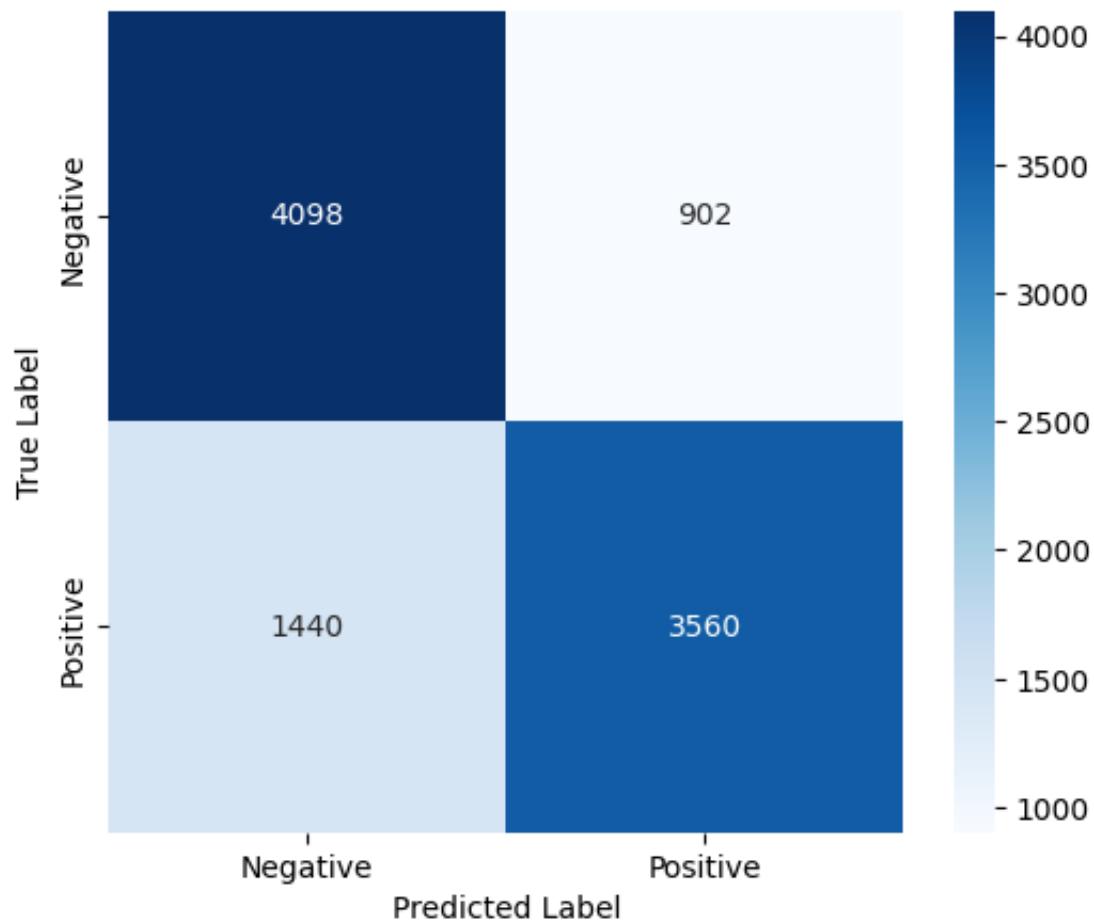
Inference Predictions for Best Configuration:

	y_true	y_pred
0	0	0
1	0	0
2	1	0
3	0	0
4	0	0
...
9995	0	1
9996	0	1
9997	0	0
9998	1	1
9999	0	0

```
9999      0      0
```

```
[10000 rows x 2 columns]
```

Confusion Matrix - ALBERT w/ LoRA on Sentiment140



◀ BITFIT

```
""" Importing BitFit packages """
```

```
import gc
import torch
import time
import pandas as pd
from tqdm import tqdm
from transformers import AutoModelForSequenceClassification, AutoTokenizer, DataCollatorWithPadding
from peft import get_peft_model, LoraConfig, TaskType
from sklearn.metrics import classification_report, f1_score
from torch.utils.data import DataLoader
```

```
""" BitFit parameter setup """
```

```
learning_rates = [5e-5, 1e-4]
batch_sizes = [8, 16]
epochs = 6
```

```
""" Training on ALBERT model using BitFit and output dataset generation (saved as
```

```
results = []
```

```
for lr in learning_rates:
    for batch_size in batch_sizes:
        print(f"Running BitFit with LR={lr}, batch_size={batch_size}")

        # loading ALBERT model
        model_name = "albert-base-v2"
        tokenizer = AutoTokenizer.from_pretrained(model_name)
        model = AutoModelForSequenceClassification.from_pretrained(model_name, num_labels=2)

        # BitFit param update config
        for name, param in model.named_parameters():
            if "bias" in name:
                param.requires_grad = True
            else:
                param.requires_grad = False

        device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
        model.to(device)

        # instantiate dataloader
        data_collator = DataCollatorWithPadding(tokenizer)
        train_dataloader = DataLoader(tokenized_dataset["train"], batch_size=batch_size, shuffle=True, collate_fn=data_collator)
```

```
test_dataloader = DataLoader(tokenized_dataset["test"], batch_size=batch_size)

# adam optimizer
optimizer = torch.optim.AdamW(filter(lambda p: p.requires_grad, model.parameters()))

# begin training
model.train()
start_time = time.time()
epoch_logs = []

for epoch in range(1, epochs + 1):
    running_loss = 0.0
    correct = 0
    total = 0
    loop = tqdm(train_dataloader, leave=True, dynamic_ncols=True, desc=f"Epoch {epoch}/{epochs}")
    for step, batch in enumerate(loop):
        batch = {
            "input_ids": batch["input_ids"].to(device),
            "attention_mask": batch["attention_mask"].to(device),
            "labels": batch["sentiment"].to(device)
        }
        with autocast(device_type='cuda'):
            outputs = model(**batch)
            loss = outputs.loss
            preds = torch.argmax(outputs.logits, dim=1)
            correct += (preds == batch['labels']).sum().item()
            total += batch['labels'].size(0)

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            running_loss += loss.item()
            running_loss += loss.item()

    avg_train_loss = running_loss / (step + 1)
    train_accuracy = correct / total

    # perform per epoch evaluation
    model.eval()
    val_running_loss = 0.0
    y_true, y_pred = [], []
    inference_start = time.time()
    with torch.no_grad():
        with autocast(device_type='cuda'):
```

```
for batch in test_dataloader:
    batch = {
        "input_ids": batch["input_ids"].to(device),
        "attention_mask": batch["attention_mask"].to(device),
        "labels": batch["sentiment"].to(device)
    }
    outputs = model(**batch)
    preds = torch.argmax(outputs.logits, dim=1)
    y_true.extend(batch["labels"].cpu().numpy())
    y_pred.extend(preds.cpu().numpy())
    val_running_loss += outputs.loss.item()

avg_val_loss = val_running_loss / len(test_dataloader)

inference_time = time.time() - inference_start

report = classification_report(y_true, y_pred, output_dict=True)
val_accuracy = report["accuracy"]
val_f1 = report["weighted avg"]["f1-score"]

epoch_logs.append({
    "epoch": epoch,
    "lr": lr,
    "batch_size": batch_size,
    "train_loss": avg_train_loss,
    "train_accuracy": train_accuracy,
    "val_loss": avg_val_loss,
    "val_accuracy": val_accuracy
})

if epoch == epochs:
    total_correct = sum(yt == yp for yt, yp in zip(y_true, y_pred))
    total_samples = len(y_true)
    accuracy = total_correct / total_samples
    f1_macro = f1_score(y_true, y_pred, average="macro")
    f1_weighted = f1_score(y_true, y_pred, average="weighted")

    print(f"\n[Final Epoch {epoch}] Inference Metrics:")
    print(f"Test Accuracy      : {accuracy:.4f}")
    print(f"F1 Score (macro)  : {f1_macro:.4f}")
    print(f"F1 Score (weighted): {f1_weighted:.4f}")
    print(f"Inference Time     : {inference_time:.2f} seconds")
    print("\nClassification Report: ALBERT w/ BitFit on Sentiment140\\")
    print(classification_report(y_true, y_pred, target_names=["Negative", "Positive"]))

    with open("sentiment140_bitfit.json", "w") as f:
        f.write(json.dumps(epoch_logs))
```

```
model.train()

end_time = time.time()
training_time = end_time - start_time

# begin datalogging per lr/bs
epoch_logs_df = pd.DataFrame(epoch_logs)
epoch_logs_df.to_csv(f"sent_albert_bitfit_epoch_logs_lr{lr}_bs{batch_size}_")

# saver inference metrics per lr/bs
metrics_summary_df = pd.DataFrame(report).transpose()
metrics_summary_df.to_csv(f"sent_albert_bitfit_inference_metrics_summary_")

# save inference predictions for the final epoch
predictions_df = pd.DataFrame({
    "y_true": y_true,
    "y_pred": y_pred
})
predictions_df.to_csv(f"sent_albert_bitfit_inference_predictions_lr{lr}_b")

# log memory usage
max_memory = torch.cuda.max_memory_allocated() / (1024 ** 3) if torch.cuda.

# save model params and metrics
results.append({
    "method": "BitFit",
    "learning_rate": lr,
    "batch_size": batch_size,
    "accuracy": val_accuracy,
    "f1": val_f1,
    "training_time": training_time,
    "inference_time": inference_time,
    "max_memory": max_memory
})

# empty cache to conserve compute
del model, tokenizer, optimizer
torch.cuda.empty_cache()
gc.collect()

# ranked performance by val acc
results = sorted(results, key=lambda x: x["accuracy"], reverse=True)

# save overall results
results_df = pd.DataFrame(results)
```

```

results_df.to_csv("sent_albert_bitfit_results.csv", index=False)

# save best final config and metrics
final_summary_df = pd.DataFrame({
    "Method": ["BitFit"],
    "Best LR": [results[0]["learning_rate"]],
    "Best Batch Size": [results[0]["batch_size"]],
    "Accuracy": [results[0]["accuracy"]],
    "F1 Score": [results[0]["f1"]],
    "Training Time (s)": [results[0]["training_time"]],
    "Inference Time (s)": [results[0]["inference_time"]],
    "Max GPU Memory (GB)": [results[0]["max_memory"]]
})
final_summary_df.to_csv("sent_albert_bf_final_comparison_bitfit.csv", index=False

print("All BitFit Grid Search Results:")
for r in results:
    print(r)

print("\nBest BitFit Configuration:")
print(results[0])

```

→ Test Accuracy : 0.7880
 F1 Score (macro) : 0.7880
 F1 Score (weighted): 0.7880
 Inference Time : 9.87 seconds

Classification Report: ALBERT w/ BitFit on Sentiment140

	precision	recall	f1-score	support
Negative	0.78	0.80	0.79	5000
Positive	0.80	0.77	0.79	5000
accuracy			0.79	10000
macro avg	0.79	0.79	0.79	10000
weighted avg	0.79	0.79	0.79	10000

Running BitFit with LR=0.0001, batch_size=8

Some weights of AlbertForSequenceClassification were not initialized from the
 You should probably TRAIN this model on a down-stream task to be able to use :

[Final Epoch 6] Inference Metrics:
 Test Accuracy : 0.7955
 F1 Score (macro) : 0.7944
 F1 Score (weighted): 0.7944
 Inference Time : 18.56 seconds

Classification Report: ALBERT w/ BitFit on Sentiment140

	precision	recall	f1-score	support
Negative	0.85	0.72	0.78	5000
Positive	0.76	0.87	0.81	5000
accuracy			0.80	10000
macro avg	0.80	0.80	0.79	10000
weighted avg	0.80	0.80	0.79	10000

Running BitFit with LR=0.0001, batch_size=16

Some weights of AlbertForSequenceClassification were not initialized from the You should probably TRAIN this model on a down-stream task to be able to use :

[Final Epoch 6] Inference Metrics:

Test Accuracy : 0.7995
 F1 Score (macro) : 0.7995
 F1 Score (weighted): 0.7995
 Inference Time : 9.84 seconds

Classification Report: ALBERT w/ BitFit on Sentiment140

	precision	recall	f1-score	support
Negative	0.80	0.79	0.80	5000
Positive	0.80	0.80	0.80	5000
accuracy			0.80	10000
macro avg	0.80	0.80	0.80	10000
weighted avg	0.80	0.80	0.80	10000

All BitFit Grid Search Results:

{'method': 'BitFit', 'learning_rate': 0.0001, 'batch_size': 16, 'accuracy': 0.7995}

```
bf_best_lr = results[0]["learning_rate"]
bf_best_bs = results[0]["batch_size"]
```

```
# Construct filename
best_report_file = f"sent_albert_bitfit_inference_metrics_summary_lr{bf_best_lr}_bs{bf_best_bs}.csv"

# Load the saved best report
best_report_df = pd.read_csv(best_report_file)
print("\nClassification Report for Best Configuration:")
print(best_report_df)
```

```
best_preds_df = pd.read_csv(f"sent_albert_bitfit_inference_predictions_lr{bf_best}")
print("\nInference Predictions for Best Configuration:")
print(best_preds_df)

y_true = best_preds_df["y_true"]
y_pred = best_preds_df["y_pred"]

cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=["Negative", "Positive"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - ALBERT w/ BitFit on Sentiment140")
plt.show()
```

[4]

Classification Report for Best Configuration:

	Unnamed: 0	precision	recall	f1-score	support
0	0	0.802342	0.7948	0.798553	5000.0000
1	1	0.796711	0.8042	0.800438	5000.0000
2	accuracy	0.799500	0.7995	0.799500	0.7995
3	macro avg	0.799526	0.7995	0.799496	10000.0000
4	weighted avg	0.799526	0.7995	0.799496	10000.0000

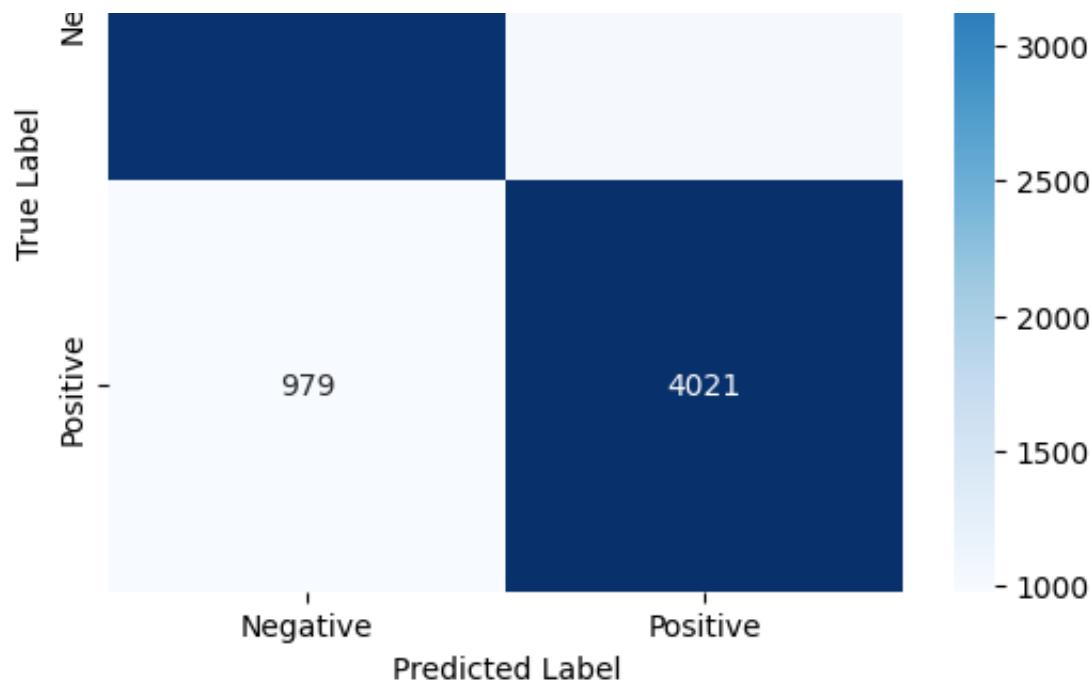
Inference Predictions for Best Configuration:

	y_true	y_pred
0	0	0
1	0	0
2	1	1
3	0	0
4	0	0
...
9995	0	0
9996	0	1
9997	0	0
9998	1	1
9999	0	0

[10000 rows x 2 columns]

Confusion Matrix - ALBERT w/ BitFit on Sentiment140





❖ Prompt Tuning

"" Importing prompt tuning packages from PEFT ""

```
import gc
from peft import PromptTuningConfig, PromptTuningInit, get_peft_model, TaskType
```

"" Prompt tuning parameter setup ""

```
lrs = [5e-5, 1e-4]
bs = [8, 16]
num_tokens = 20
epochs = 6
```

"" Training and evaluation loop with hyperparameter grid search ""

```
results = []
```

```
for lr in lrs:
    for batch_size in bs:
        print(f"Running Prompt Tuning with LR={lr}, batch_size={batch_size}")
```

```
# loading ALBERT model
```

```
model_name = "albert-base-v2"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForSequenceClassification.from_pretrained(model_name, num_
```



```
# prompt tuning config
peft_config = PromptTuningConfig(
    task_type=TaskType.SEQ_CLS,
    num_virtual_tokens=num_tokens,
    tokenizer_name_or_path=tokenizer.name_or_path,
    prompt_tuning_init=PromptTuningInit.RANDOM,
)
prompt_model = get_peft_model(model, peft_config)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
prompt_model.to(device)
```



```
# instantiate dataloader
data_collator = DataCollatorWithPadding(tokenizer)
train_dataloader = DataLoader(tokenized_dataset["train"], batch_size=batch_
test_dataloader = DataLoader(tokenized_dataset["test"], batch_size=batch_si
```



```
# adam optimization
optimizer = torch.optim.AdamW(prompt_model.parameters(), lr=lr)
```



```
# begin training
prompt_model.train()
start_time = time.time()
epoch_logs = []
```



```
for epoch in range(1, epochs + 1):
    running_loss = 0.0
    correct = 0
    total = 0
    loop = tqdm(train_dataloader, leave=True, dynamic_ncols=True, desc=f"Epoch {epoch}/{epochs}")
    for step, batch in enumerate(loop):
        batch = {
            "input_ids": batch["input_ids"].to(device),
            "attention_mask": batch["attention_mask"].to(device),
            "labels": batch["sentiment"].to(device)
        }
        with autocast(device_type='cuda'):
            outputs = model(**batch)
            loss = outputs.loss
            preds = torch.argmax(outputs.logits, dim=1)
            correct += (preds == batch['labels']).sum().item()
            total += batch['labels'].size(0)
```

```
optimizer.zero_grad()
loss.backward()
optimizer.step()

running_loss += loss.item()

avg_train_loss = running_loss / (step + 1)
train_accuracy = correct / total

# perform per epoch evaluation
prompt_model.eval()
val_running_loss = 0.0
y_true, y_pred = [], []
with torch.no_grad():
    with autocast(device_type='cuda'):
        for batch in test_dataloader:
            batch = {
                "input_ids": batch["input_ids"].to(device),
                "attention_mask": batch["attention_mask"].to(device),
                "labels": batch["sentiment"].to(device)
            }
            outputs = model(**batch)
            preds = torch.argmax(outputs.logits, dim=1)
            y_true.extend(batch["labels"].cpu().numpy())
            y_pred.extend(preds.cpu().numpy())
            val_running_loss += outputs.loss.item()

avg_val_loss = val_running_loss / len(test_dataloader)

inference_time = time.time() - start_time

report = classification_report(y_true, y_pred, output_dict=True)
val_accuracy = report["accuracy"]
val_f1 = report["weighted avg"]["f1-score"]

# print classification report on final epoch
if epoch == epochs:
    total_correct = sum(yt == yp for yt, yp in zip(y_true, y_pred))
    total_samples = len(y_true)

    accuracy = total_correct / total_samples
    f1_macro = f1_score(y_true, y_pred, average="macro")
    f1_weighted = f1_score(y_true, y_pred, average="weighted")
```

```
        print(f"\n[Final Epoch {epoch}] Inference Metrics:")
        print(f"Test Accuracy      : {accuracy:.4f}")
        print(f"F1 Score (macro)  : {f1_macro:.4f}")
        print(f"F1 Score (weighted): {f1_weighted:.4f}")
        print(f"Inference Time    : {inference_time:.2f} seconds")
        print("\nClassification Report: - ALBERT w/ Prompt Tuning on Sentin")
        print(classification_report(y_true, y_pred, target_names=["Negative", "Positive"]))

    epoch_logs.append({
        "epoch": epoch,
        "lr": lr,
        "batch_size": batch_size,
        "train_loss": avg_train_loss,
        "train_accuracy": train_accuracy,
        "val_loss": avg_val_loss,
        "val_accuracy": val_accuracy
    })

    prompt_model.train()

end_time = time.time()
training_time = end_time - start_time

# begin datalogging per lr/bs
epoch_logs_df = pd.DataFrame(epoch_logs)
epoch_logs_df.to_csv(f"sent_albert_prompt_epoch_logs_lr{lr}_bs{batch_size}.csv")

# save inference metrics per lr/bs
metrics_summary_df = pd.DataFrame(report).transpose()
metrics_summary_df.to_csv(f"sent_albert_prompt_inference_metrics_summary_lr{lr}.csv")

# Save inference predictions for the final epoch
predictions_df = pd.DataFrame({
    "y_true": y_true,
    "y_pred": y_pred
})
predictions_df.to_csv(f"sent_albert_prompt_inference_predictions_lr{lr}_bs{batch_size}.csv")

# log memory usage
max_memory = torch.cuda.max_memory_allocated() / (1024 ** 3) if torch.cuda.is_available() else None
print(f"Max memory usage: {max_memory} GB")

# save model params and metrics
results.append({
    "method": "Prompt Tuning",
    "epoch": epoch,
    "lr": lr,
    "batch_size": batch_size,
    "train_loss": avg_train_loss,
    "train_accuracy": train_accuracy,
    "val_loss": avg_val_loss,
    "val_accuracy": val_accuracy,
    "inference_time": inference_time,
    "max_memory": max_memory
})
```

```

        "learning_rate": lr,
        "batch_size": batch_size,
        "accuracy": val_accuracy,
        "f1": val_f1,
        "training_time": training_time,
        "inference_time": inference_time,
        "max_memory": max_memory
    })

# empty cache to conserve compute
del prompt_model, model, tokenizer, optimizer
torch.cuda.empty_cache()
gc.collect()

# ranked performance by val acc
results = sorted(results, key=lambda x: x["accuracy"], reverse=True)

# save overall results
results_df = pd.DataFrame(results)
results_df.to_csv("sent_albert_prompt_results.csv", index=False)

# save best final config and metrics
final_summary_df = pd.DataFrame({
    "Method": ["Prompt Tuning"],
    "Best LR": [results[0]["learning_rate"]],
    "Best Batch Size": [results[0]["batch_size"]],
    "Accuracy": [results[0]["accuracy"]],
    "F1 Score": [results[0]["f1"]],
    "Training Time (s)": [results[0]["training_time"]],
    "Max GPU Memory (GB)": [results[0]["max_memory"]]
})
final_summary_df.to_csv("sent_albert_prompt_final_comparison_prompt_tuning.csv", index=False)

print("All Prompt Tuning Grid Search Results:")
for r in results:
    print(r)

print("\nBest Configuration:")
print(results[0])

```

	precision	recall	f1-score	support
Negative	0.67	0.69	0.68	5000
Positive	0.68	0.66	0.67	5000

accuracy			0.68	10000
macro avg	0.68	0.68	0.68	10000
weighted avg	0.68	0.68	0.68	10000

Running Prompt Tuning with LR=0.0001, batch_size=8

Some weights of AlbertForSequenceClassification were not initialized from the You should probably TRAIN this model on a down-stream task to be able to use :

Epoch 1/6: 100%		5000/5000	[01:26<00:00, 57.87it/s]
Epoch 2/6: 100%		5000/5000	[01:25<00:00, 58.16it/s]
Epoch 3/6: 100%		5000/5000	[01:26<00:00, 57.72it/s]
Epoch 4/6: 100%		5000/5000	[01:28<00:00, 56.62it/s]
Epoch 5/6: 100%		5000/5000	[01:27<00:00, 57.04it/s]
Epoch 6/6: 100%		5000/5000	[01:27<00:00, 57.46it/s]

[Final Epoch 6] Inference Metrics:

Test Accuracy : 0.6941
 F1 Score (macro) : 0.6937
 F1 Score (weighted): 0.6937
 Inference Time : 640.50 seconds

Classification Report: - ALBERT w/ Prompt Tuning on Sentiment140

	precision	recall	f1-score	support
Negative	0.71	0.66	0.68	5000
Positive	0.68	0.73	0.71	5000
accuracy			0.69	10000
macro avg	0.70	0.69	0.69	10000
weighted avg	0.70	0.69	0.69	10000

Running Prompt Tuning with LR=0.0001, batch_size=16

Some weights of AlbertForSequenceClassification were not initialized from the You should probably TRAIN this model on a down-stream task to be able to use :

Epoch 1/6: 100%		2500/2500	[00:45<00:00, 54.46it/s]
Epoch 2/6: 100%		2500/2500	[00:45<00:00, 54.68it/s]
Epoch 3/6: 100%		2500/2500	[00:45<00:00, 54.51it/s]
Epoch 4/6: 100%		2500/2500	[00:46<00:00, 54.31it/s]
Epoch 5/6: 100%		2500/2500	[00:45<00:00, 54.39it/s]
Epoch 6/6: 100%		2500/2500	[00:45<00:00, 54.93it/s]

[Final Epoch 6] Inference Metrics:

Test Accuracy : 0.6878
 F1 Score (macro) : 0.6878
 F1 Score (weighted): 0.6878
 Inference Time : 337.55 seconds

Classification Report: - ALBERT w/ Prompt Tuning on Sentiment140

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Negative	0.69	0.69	0.69	5000
Positive	0.69	0.69	0.69	5000

```
prompt_best_lr = results[0]["learning_rate"]
prompt_best_bs = results[0]["batch_size"]

# Construct filename
best_report_file = f"sent_albert_prompt_inference_metrics_summary_lr{prompt_best_"

# Load the saved best report
best_report_df = pd.read_csv(best_report_file)
print("\nClassification Report for Best Configuration:")
print(best_report_df)

best_preds_df = pd.read_csv(f"sent_albert_prompt_inference_predictions_lr{prompt_"
print("\nInference Predictions for Best Configuration:")
print(best_preds_df)

y_true = best_preds_df["y_true"]
y_pred = best_preds_df["y_pred"]

cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=["Negative", "Positive"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - ALBERT w/ Prompt Tuning on Sentiment140\n")
plt.show()
```



Classification Report for Best Configuration:

	Unnamed: 0	precision	recall	f1-score	support
0	0	0.710019	0.6562	0.68205	5000.0000
1	1	0.680424	0.7320	0.70527	5000.0000
2	accuracy	0.694100	0.6941	0.69410	0.6941
3	macro avg	0.695222	0.6941	0.69366	10000.0000
4	weighted avg	0.695222	0.6941	0.69366	10000.0000

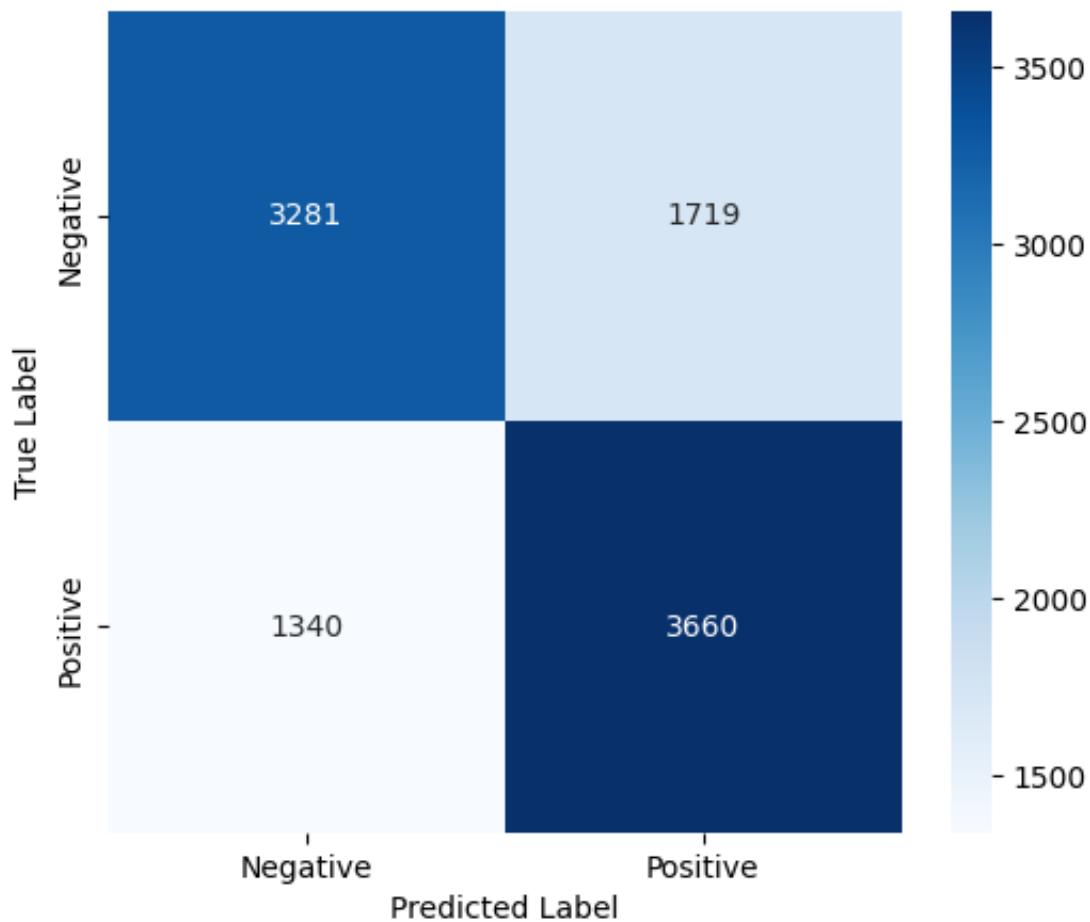
Inference Predictions for Best Configuration:

	y_true	y_pred
0	0	0
1	0	1
2	1	0
3	0	0

```
4          0          0
...
9995      0          1
9996      0          1
9997      0          1
9998      1          1
9999      0          0
```

[10000 rows x 2 columns]

Confusion Matrix - ALBERT w/ Prompt Tuning on Sentiment140



✓ Begin Visualization of Sentiment140 Results

Load all dataframes from above training of 3 PEFT methods

"""\ Output from our ALBERT_Sentiment140.ipynb file, we read in the .csv files here. Note that these .csv file paths suggest they should be uploaded to session storage.

```
# ALBERT PEFT-wise results across bf/lora/prompt tuning
albert_bf_results = pd.read_csv('/content/sent_albert_bitfit_results.csv')
albert_lora_results = pd.read_csv('/content/sent_albert_lora_results.csv')
albert_prompt_results = pd.read_csv('/content/sent_albert_prompt_results.csv')

# ALBERT per-epoch performance logs (for LC generation) across bf/lora/prompt tuning
albert_lora_epochs_lr5_bs8 = pd.read_csv('/content/sent_albert_lora_epoch_logs_lr5_bs8.csv')
albert_lora_epochs_lr5_bs16 = pd.read_csv('/content/sent_albert_lora_epoch_logs_lr5_bs16.csv')
albert_lora_epochs_lr1_bs8 = pd.read_csv('/content/sent_albert_lora_epoch_logs_lr1_bs8.csv')
albert_lora_epochs_lr1_bs16 = pd.read_csv('/content/sent_albert_lora_epoch_logs_lr1_bs16.csv')
albert_bf_epochs_lr5_bs8 = pd.read_csv('/content/sent_albert_bitfit_epoch_logs_lr5_bs8.csv')
albert_bf_epochs_lr5_bs16 = pd.read_csv('/content/sent_albert_bitfit_epoch_logs_lr5_bs16.csv')
albert_bf_epochs_lr1_bs8 = pd.read_csv('/content/sent_albert_bitfit_epoch_logs_lr1_bs8.csv')
albert_bf_epochs_lr1_bs16 = pd.read_csv('/content/sent_albert_bitfit_epoch_logs_lr1_bs16.csv')
albert_prompt_epochs_lr5_bs8 = pd.read_csv('/content/sent_albert_prompt_epoch_logs_lr5_bs8.csv')
albert_prompt_epochs_lr5_bs16 = pd.read_csv('/content/sent_albert_prompt_epoch_logs_lr5_bs16.csv')
albert_prompt_epochs_lr1_bs8 = pd.read_csv('/content/sent_albert_prompt_epoch_logs_lr1_bs8.csv')
albert_prompt_epochs_lr1_bs16 = pd.read_csv('/content/sent_albert_prompt_epoch_logs_lr1_bs16.csv')

# ALBERT inference performance metric summary across bf/lora/prompt tuning
albert_bf_inf_lr5_bs8 = pd.read_csv('/content/sent_albert_bitfit_inference_metric_lr5_bs8.csv')
albert_bf_inf_lr5_bs16 = pd.read_csv('/content/sent_albert_bitfit_inference_metric_lr5_bs16.csv')
albert_bf_inf_lr1_bs8 = pd.read_csv('/content/sent_albert_bitfit_inference_metric_lr1_bs8.csv')
albert_bf_inf_lr1_bs16 = pd.read_csv('/content/sent_albert_bitfit_inference_metric_lr1_bs16.csv')
albert_lora_inf_lr5_bs8 = pd.read_csv('/content/sent_albert_lora_inference_metric_lr5_bs8.csv')
albert_lora_inf_lr5_bs16 = pd.read_csv('/content/sent_albert_lora_inference_metric_lr5_bs16.csv')
albert_lora_inf_lr1_bs8 = pd.read_csv('/content/sent_albert_lora_inference_metric_lr1_bs8.csv')
albert_lora_inf_lr1_bs16 = pd.read_csv('/content/sent_albert_lora_inference_metric_lr1_bs16.csv')
albert_prompt_inf_lr5_bs8 = pd.read_csv('/content/sent_albert_prompt_inference_metric_lr5_bs8.csv')
albert_prompt_inf_lr5_bs16 = pd.read_csv('/content/sent_albert_prompt_inference_metric_lr5_bs16.csv')
albert_prompt_inf_lr1_bs8 = pd.read_csv('/content/sent_albert_prompt_inference_metric_lr1_bs8.csv')
albert_prompt_inf_lr1_bs16 = pd.read_csv('/content/sent_albert_prompt_inference_metric_lr1_bs16.csv')

# ALBERT inference predictions across bf/lora/prompt tuning
albert_bf_preds_lr5_bs8 = pd.read_csv('/content/sent_albert_bitfit_inference_pred_lr5_bs8.csv')
albert_bf_preds_lr5_bs16 = pd.read_csv('/content/sent_albert_bitfit_inference_pred_lr5_bs16.csv')
albert_bf_preds_lr1_bs8 = pd.read_csv('/content/sent_albert_bitfit_inference_pred_lr1_bs8.csv')
albert_bf_preds_lr1_bs16 = pd.read_csv('/content/sent_albert_bitfit_inference_pred_lr1_bs16.csv')
albert_lora_preds_lr5_bs8 = pd.read_csv('/content/sent_albert_lora_inference_pred_lr5_bs8.csv')
albert_lora_preds_lr5_bs16 = pd.read_csv('/content/sent_albert_lora_inference_pred_lr5_bs16.csv')
albert_lora_preds_lr1_bs8 = pd.read_csv('/content/sent_albert_lora_inference_pred_lr1_bs8.csv')
albert_lora_preds_lr1_bs16 = pd.read_csv('/content/sent_albert_lora_inference_pred_lr1_bs16.csv')
albert_prompt_preds_lr5_bs8 = pd.read_csv('/content/sent_albert_prompt_inference_pred_lr5_bs8.csv')
albert_prompt_preds_lr5_bs16 = pd.read_csv('/content/sent_albert_prompt_inference_pred_lr5_bs16.csv')
albert_prompt_preds_lr1_bs8 = pd.read_csv('/content/sent_albert_prompt_inference_pred_lr1_bs8.csv')
albert_prompt_preds_lr1_bs16 = pd.read_csv('/content/sent_albert_prompt_inference_pred_lr1_bs16.csv')
```

```

albert_prompt_preds_lr1_bs8 = pd.read_csv('/content/sent_albert_prompt_inference_
albert_prompt_preds_lr1_bs16 = pd.read_csv('/content/sent_albert_prompt_inference.

# ALBERT PEFT method intra-comparison based on hyperparameter settings, per bf/lo
albert_bf_final_comparison = pd.read_csv('/content/sent_albert_bf_final_compariso
albert_lora_final_comparison = pd.read_csv('/content/sent_albert_lora_final_compa
albert_prompt_final_comparison = pd.read_csv('/content/sent_albert_prompt_final_c

```

BitFit Learning Curves

```

# All BitFit Train/Val Acc Learning Curve
plt.figure(figsize=(10,5))
sns.lineplot(data=albert_bf_epochs_lr5_bs8, x="epoch", y="train_accuracy", label=
sns.lineplot(data=albert_bf_epochs_lr5_bs8, x="epoch", y="val_accuracy", label="V
sns.lineplot(data=albert_bf_epochs_lr5_bs16, x="epoch", y="train_accuracy", label:
sns.lineplot(data=albert_bf_epochs_lr5_bs16, x="epoch", y="val_accuracy", label=""
sns.lineplot(data=albert_bf_epochs_lr1_bs8, x="epoch", y="train_accuracy", label=
sns.lineplot(data=albert_bf_epochs_lr1_bs8, x="epoch", y="val_accuracy", label="V
sns.lineplot(data=albert_bf_epochs_lr1_bs16, x="epoch", y="train_accuracy", label:
sns.lineplot(data=albert_bf_epochs_lr1_bs16, x="epoch", y="val_accuracy", label=""
plt.title("Learning Curve: Accuracy - ALBERT w/ BitFit on Sentiment140")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.grid(True)
plt.show()

```

```

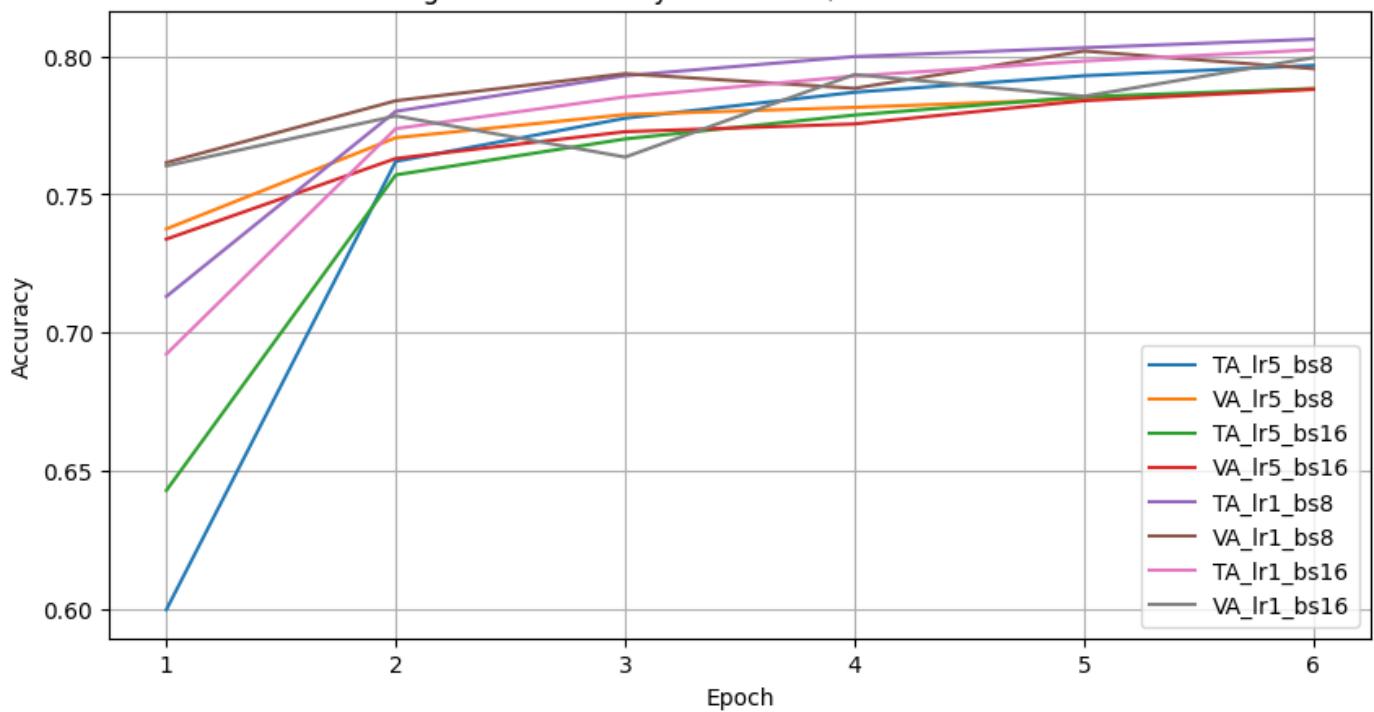
# All BitFit Training and Validation Loss
plt.figure(figsize=(10,5))
sns.lineplot(data=albert_bf_epochs_lr5_bs8, x="epoch", y="train_loss", label="TL_
sns.lineplot(data=albert_bf_epochs_lr5_bs8, x="epoch", y="val_loss", label="VL_lr
sns.lineplot(data=albert_bf_epochs_lr5_bs16, x="epoch", y="train_loss", label="TL_
sns.lineplot(data=albert_bf_epochs_lr5_bs16, x="epoch", y="val_loss", label="VL_l
sns.lineplot(data=albert_bf_epochs_lr1_bs8, x="epoch", y="train_loss", label="TL_
sns.lineplot(data=albert_bf_epochs_lr1_bs8, x="epoch", y="val_loss", label="VL_lr
sns.lineplot(data=albert_bf_epochs_lr1_bs16, x="epoch", y="train_loss", label="TL_
sns.lineplot(data=albert_bf_epochs_lr1_bs16, x="epoch", y="val_loss", label="VL_l
plt.title("Learning Curve: Loss - ALBERT w/ BitFit on Sentiment140")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)

```

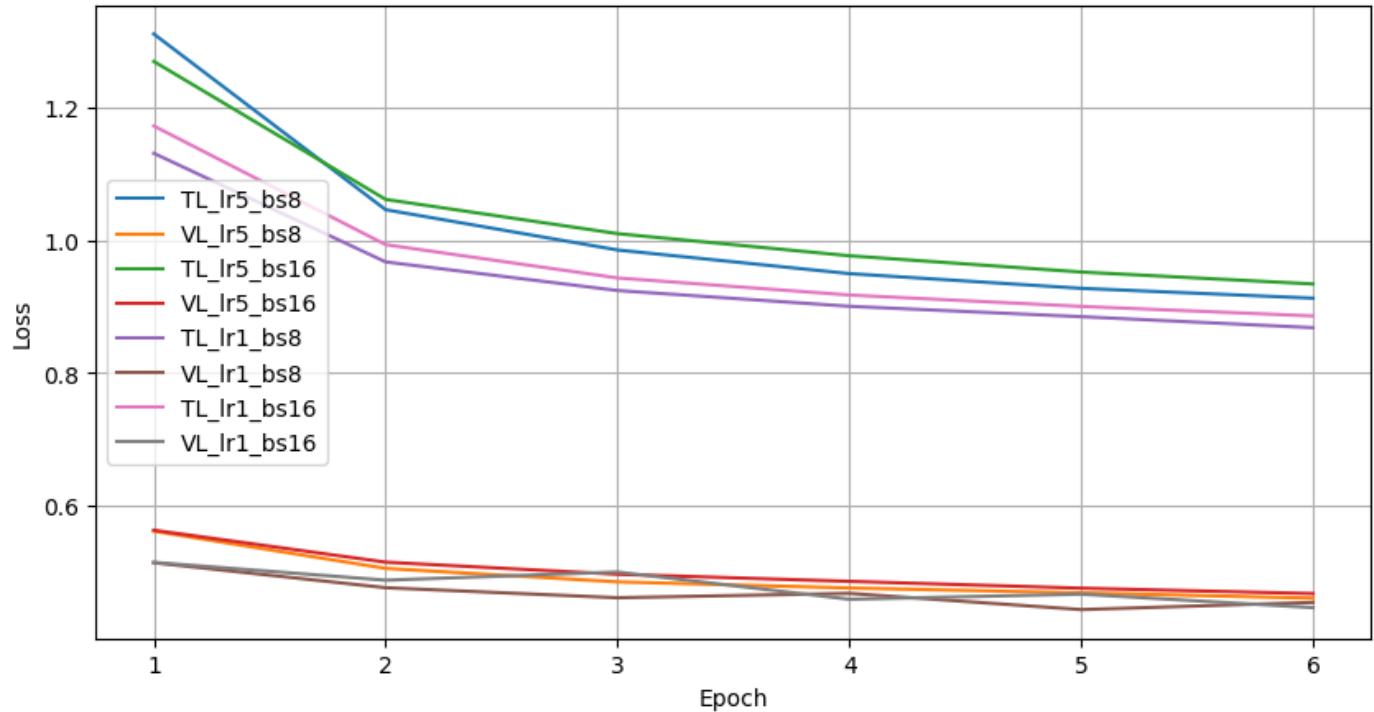
```
plt.show()
```



Learning Curve: Accuracy - ALBERT w/ BitFit on Sentiment140



Learning Curve: Loss - ALBERT w/ BitFit on Sentiment140



```
# Best BitFit Train/Val Acc Learning Curve

albert_bf_epochs_map = {
    (5, 8): albert_bf_epochs_lr5_bs8,
    (5, 16): albert_bf_epochs_lr5_bs16,
    (1, 8): albert_bf_epochs_lr1_bs8,
    (1, 16): albert_bf_epochs_lr1_bs16
}

bf_lr_mapping = {
    5e-5: 5,
    1e-4: 1
}

bf_best_lr_tag = bf_lr_mapping[bf_best_lr]
bf_best_bs_tag = bf_best_bs

bf_epochs = albert_bf_epochs_map[(bf_best_lr_tag, bf_best_bs_tag)]

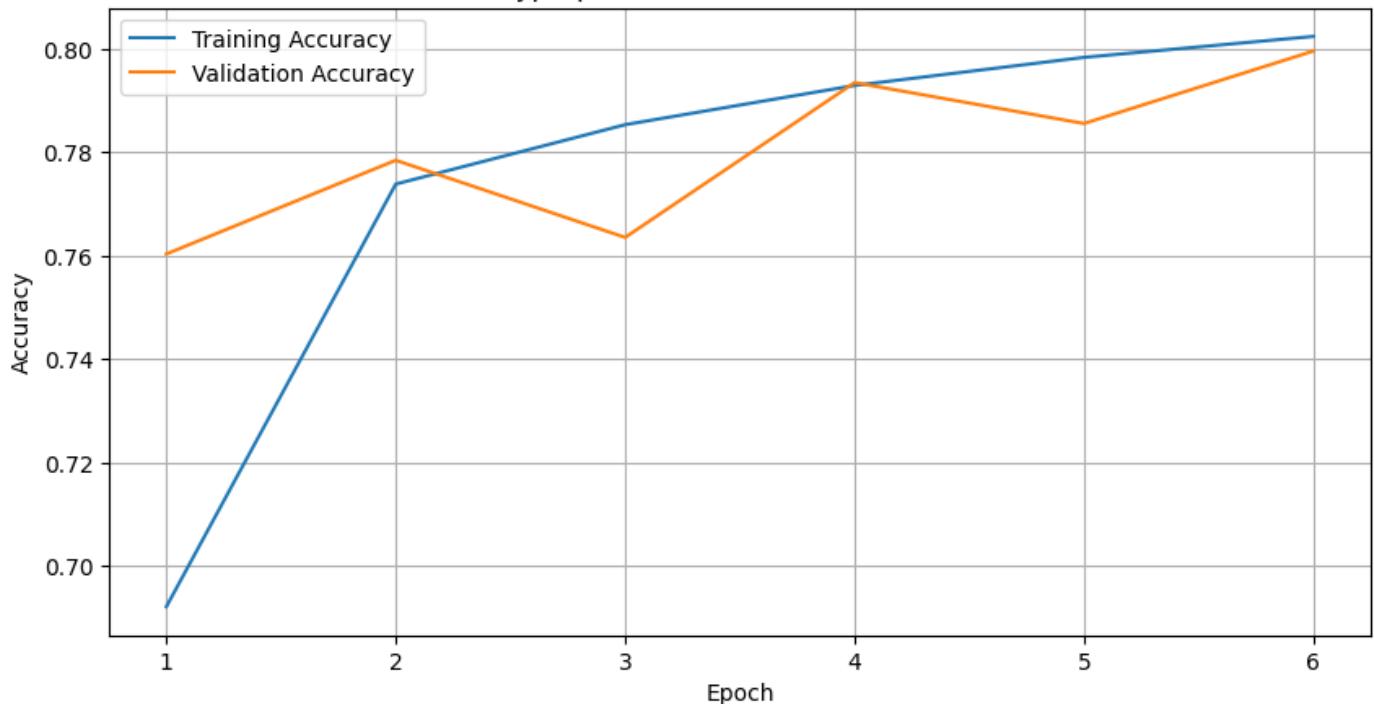
# Best BitFit Training and Validation Accuracy
plt.figure(figsize=(10,5))
sns.lineplot(data=bf_epochs, x="epoch", y="train_accuracy", label="Training Accuracy")
sns.lineplot(data=bf_epochs, x="epoch", y="val_accuracy", label="Validation Accuracy")
plt.title(f"Learning Curve: Best Accuracy - ALBERT w/ BitFit on Sentiment140\nBest Hyperparameters: {bf_best_lr_tag} lr, {bf_best_bs_tag} bs")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.grid(True)
plt.show()

# Best BitFit Training and Validation Loss
plt.figure(figsize=(10,5))
sns.lineplot(data=bf_epochs, x="epoch", y="train_loss", label="Training Loss")
sns.lineplot(data=bf_epochs, x="epoch", y="val_loss", label="Validation Loss")
plt.title(f"Learning Curve: Best Loss - ALBERT w/ BitFit on Sentiment140\nBest Hyperparameters: {bf_best_lr_tag} lr, {bf_best_bs_tag} bs")
plt.xlabel("Epoch")
plt.ylabel("Loss")
```

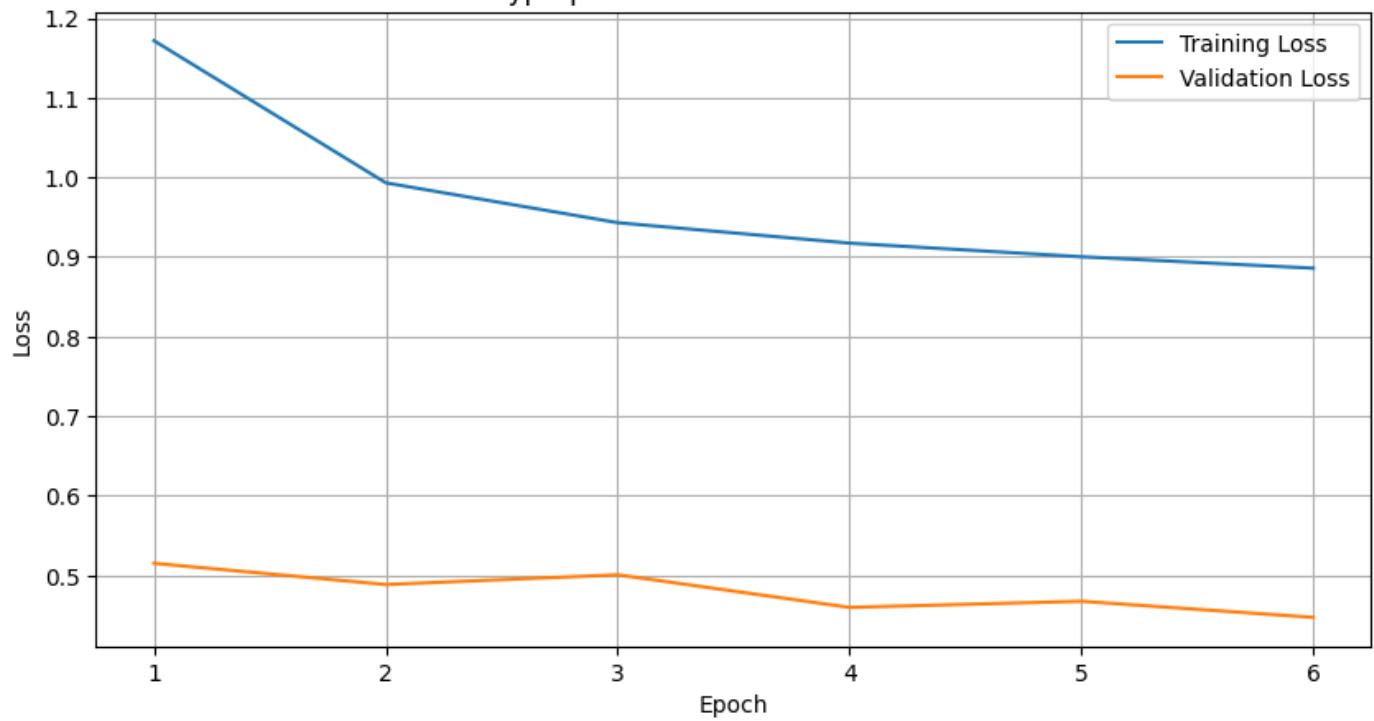
```
plt.legend()  
plt.grid(True)  
plt.show()
```



Learning Curve: Best Accuracy - ALBERT w/ BitFit on Sentiment140
Best Hyperparameters: LR: 0.0001 and BS: 16



Learning Curve: Best Loss - ALBERT w/ BitFit on Sentiment140
Best Hyperparameters: LR: 0.0001 and BS: 16

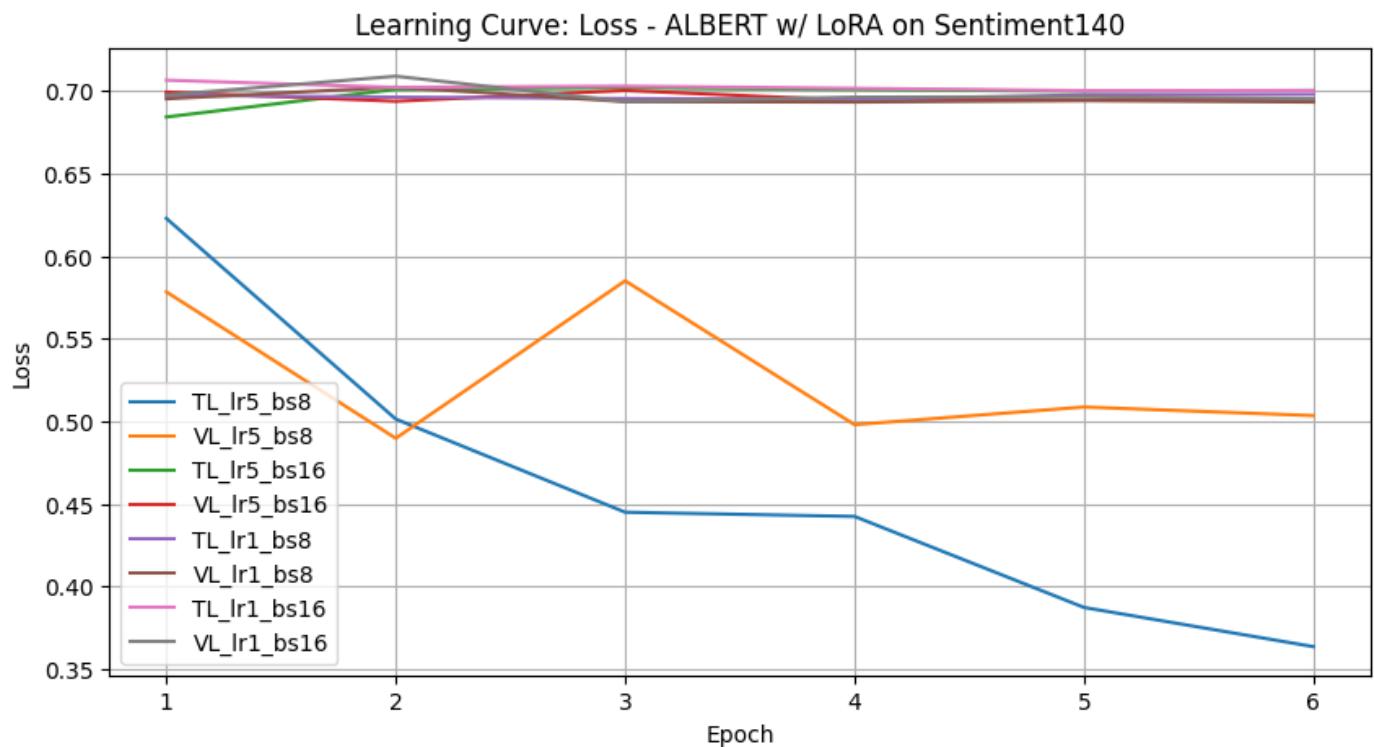
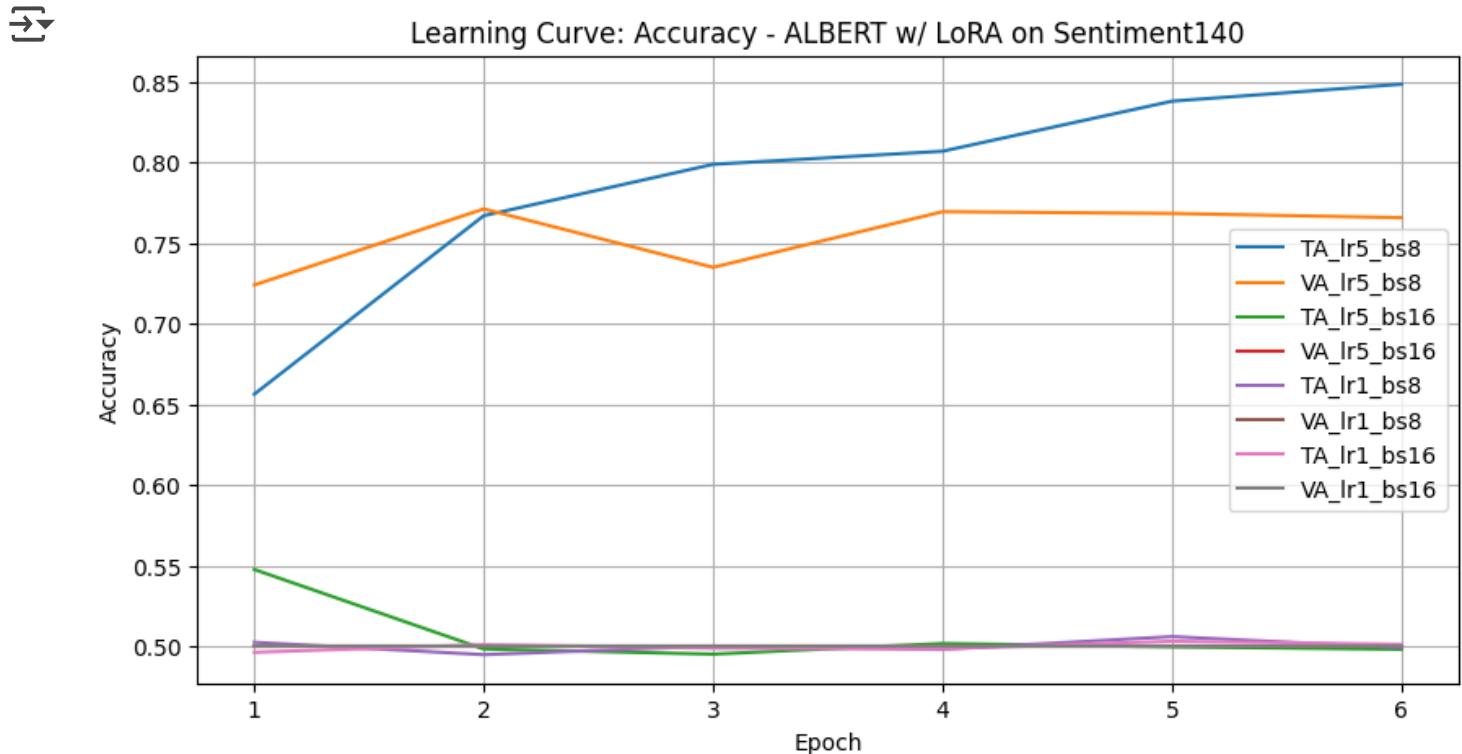


LoRA Learning Curves

```
# All LoRA Train/Val Acc Learning Curve
plt.figure(figsize=(10,5))
sns.lineplot(data=albert_lora_epochs_lr5_bs8, x="epoch", y="train_accuracy", label="Train Accuracy - lr5_bs8")
sns.lineplot(data=albert_lora_epochs_lr5_bs8, x="epoch", y="val_accuracy", label="Val Accuracy - lr5_bs8")
sns.lineplot(data=albert_lora_epochs_lr5_bs16, x="epoch", y="train_accuracy", label="Train Accuracy - lr5_bs16")
sns.lineplot(data=albert_lora_epochs_lr5_bs16, x="epoch", y="val_accuracy", label="Val Accuracy - lr5_bs16")
sns.lineplot(data=albert_lora_epochs_lr1_bs8, x="epoch", y="train_accuracy", label="Train Accuracy - lr1_bs8")
sns.lineplot(data=albert_lora_epochs_lr1_bs8, x="epoch", y="val_accuracy", label="Val Accuracy - lr1_bs8")
sns.lineplot(data=albert_lora_epochs_lr1_bs16, x="epoch", y="train_accuracy", label="Train Accuracy - lr1_bs16")
sns.lineplot(data=albert_lora_epochs_lr1_bs16, x="epoch", y="val_accuracy", label="Val Accuracy - lr1_bs16")
plt.title("Learning Curve: Accuracy - ALBERT w/ LoRA on Sentiment140")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.grid(True)
plt.show()

# All LoRA Training and Validation Loss
plt.figure(figsize=(10,5))
sns.lineplot(data=albert_lora_epochs_lr5_bs8, x="epoch", y="train_loss", label="Train Loss - lr5_bs8")
sns.lineplot(data=albert_lora_epochs_lr5_bs8, x="epoch", y="val_loss", label="Val Loss - lr5_bs8")
sns.lineplot(data=albert_lora_epochs_lr5_bs16, x="epoch", y="train_loss", label="Train Loss - lr5_bs16")
sns.lineplot(data=albert_lora_epochs_lr5_bs16, x="epoch", y="val_loss", label="Val Loss - lr5_bs16")
sns.lineplot(data=albert_lora_epochs_lr1_bs8, x="epoch", y="train_loss", label="Train Loss - lr1_bs8")
sns.lineplot(data=albert_lora_epochs_lr1_bs8, x="epoch", y="val_loss", label="Val Loss - lr1_bs8")
sns.lineplot(data=albert_lora_epochs_lr1_bs16, x="epoch", y="train_loss", label="Train Loss - lr1_bs16")
sns.lineplot(data=albert_lora_epochs_lr1_bs16, x="epoch", y="val_loss", label="Val Loss - lr1_bs16")
plt.title("Learning Curve: Loss - ALBERT w/ LoRA on Sentiment140")
plt.xlabel("Epoch")
plt.ylabel("Loss")
```

```
plt.legend()  
plt.grid(True)  
plt.show()
```



```
# Best LoRA Train/Val Acc Learning Curve

albert_lora_epochs_map = {
    (5, 8): albert_lora_epochs_lr5_bs8,
    (5, 16): albert_lora_epochs_lr5_bs16,
    (1, 8): albert_lora_epochs_lr1_bs8,
    (1, 16): albert_lora_epochs_lr1_bs16
}

lora_lr_mapping = {
    5e-5: 5,
    1e-4: 1
}

lora_best_lr_tag = lora_lr_mapping[lora_best_lr]
lora_best_bs_tag = lora_best_bs

lora_epochs = albert_lora_epochs_map[(lora_best_lr_tag, lora_best_bs_tag)]

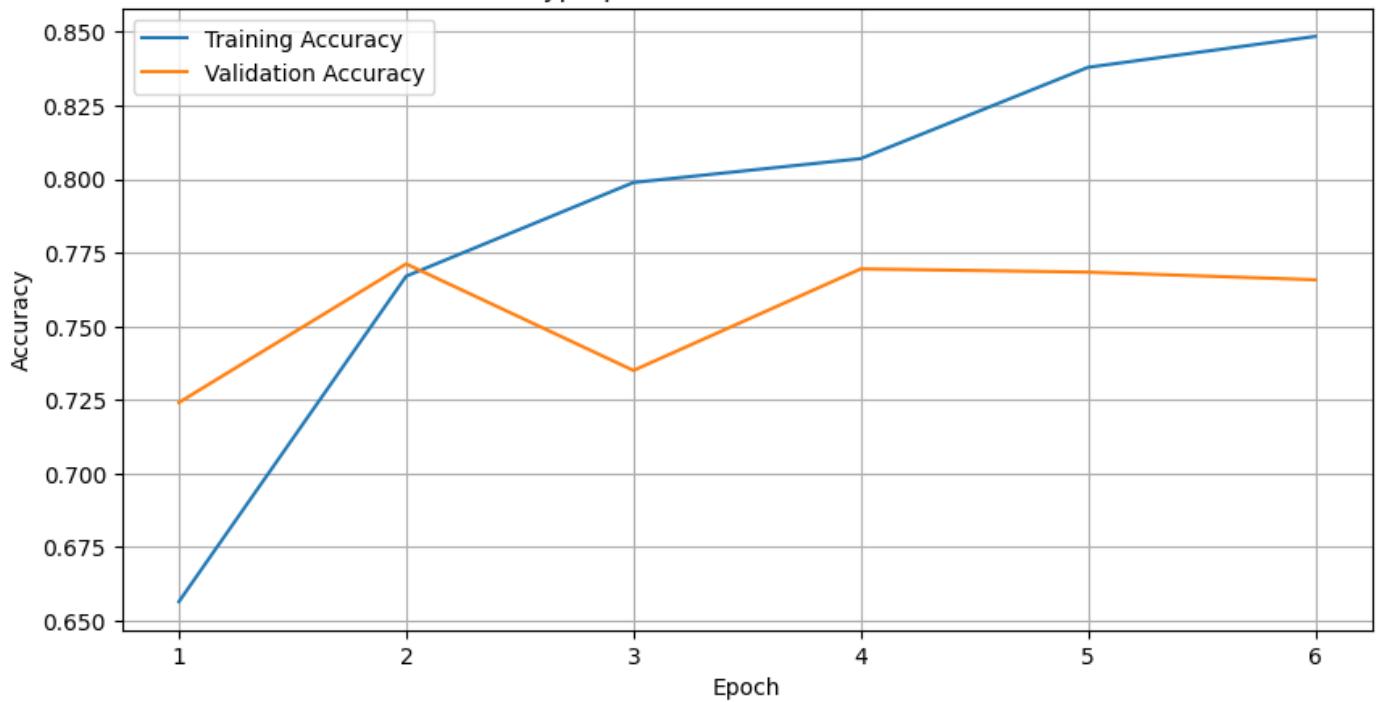
# Best LoRA Training and Validation Accuracy
plt.figure(figsize=(10,5))
sns.lineplot(data=lora_epochs, x="epoch", y="train_accuracy", label="Training Accuracy")
sns.lineplot(data=lora_epochs, x="epoch", y="val_accuracy", label="Validation Accuracy")
plt.title(f"Learning Curve: Best Accuracy - ALBERT w/ LoRA on Sentiment140\nBest Hyperparameters: {lora_best_lr_tag} and {lora_best_bs_tag}")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.grid(True)
plt.show()

# Best LoRA Training and Validation Loss
plt.figure(figsize=(10,5))
sns.lineplot(data=lora_epochs, x="epoch", y="train_loss", label="Training Loss")
sns.lineplot(data=lora_epochs, x="epoch", y="val_loss", label="Validation Loss")
plt.title(f"Learning Curve: Best Loss - ALBERT w/ LoRA on Sentiment140\nBest Hyperparameters: {lora_best_lr_tag} and {lora_best_bs_tag}")
```

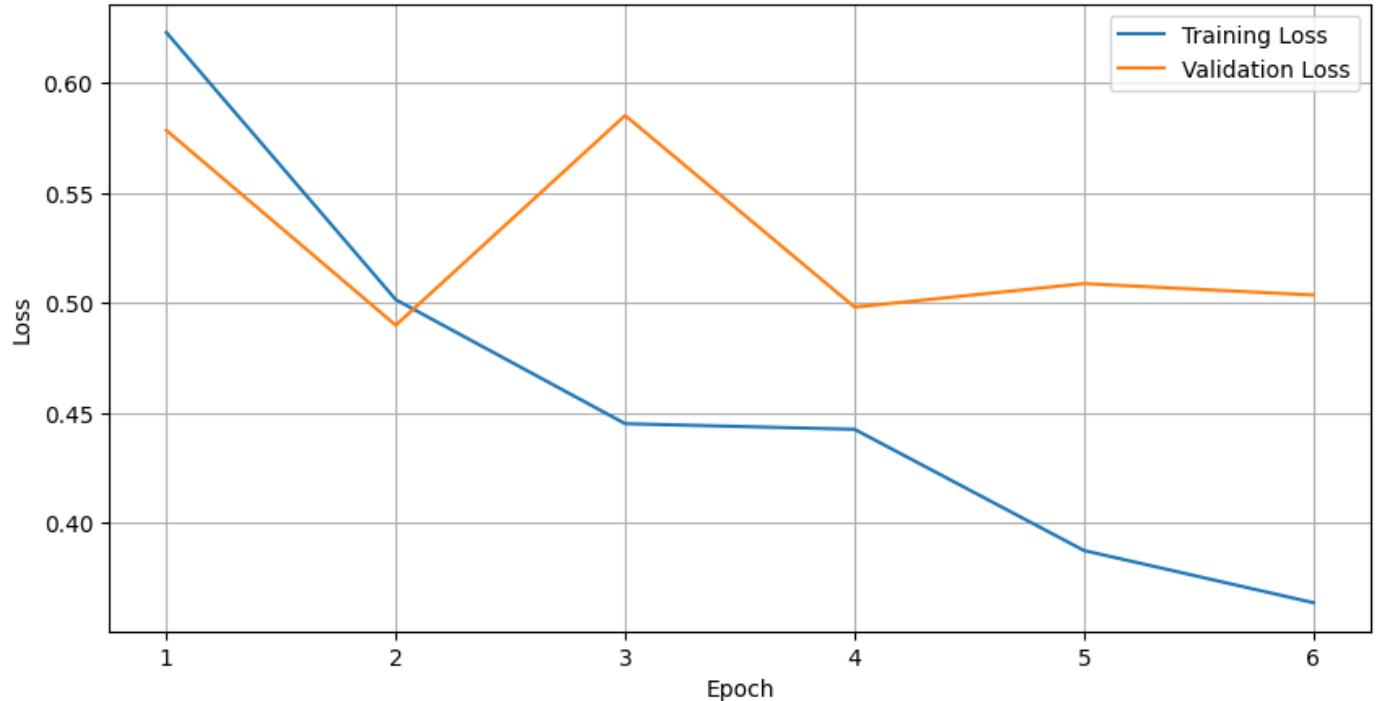
```
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()
```



Learning Curve: Best Accuracy - ALBERT w/ LoRA on Sentiment140
Best Hyperparameters: LR: 5e-05 and BS: 8



Learning Curve: Best Loss - ALBERT w/ LoRA on Sentiment140
Best Hyperparameters: LR: 5e-05 and BS: 8

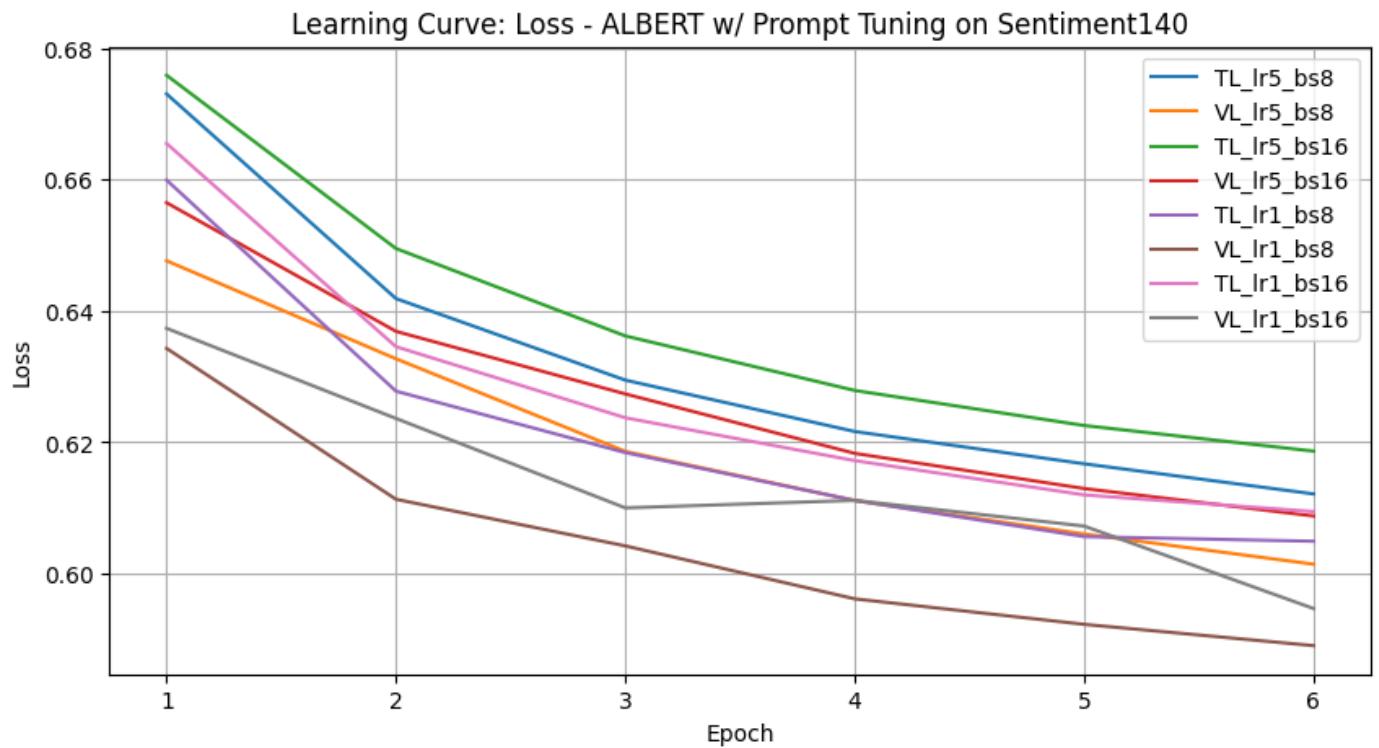
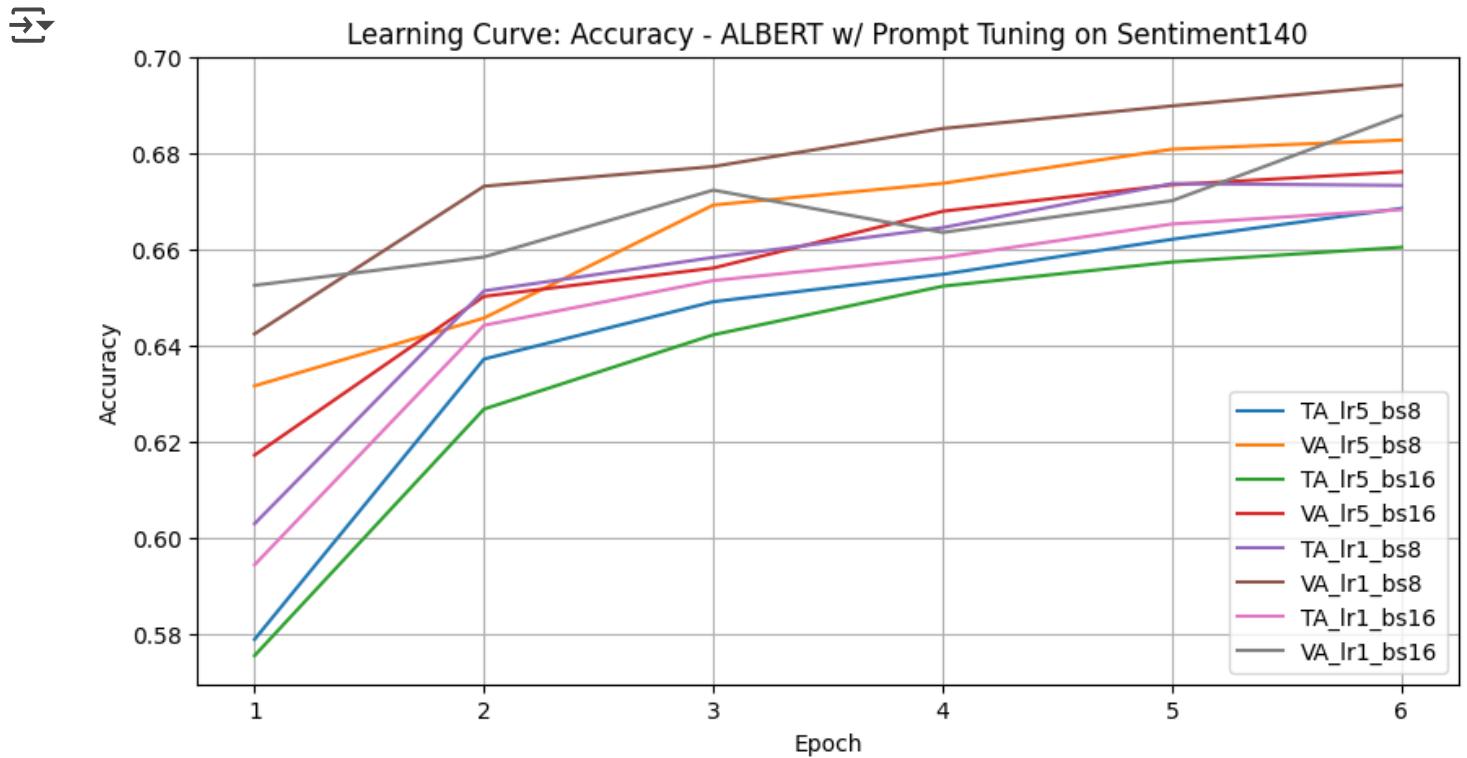


Prompt Tuning Learning Curves

```
# All Prompt Tuning Train/Val Acc Learning Curve
plt.figure(figsize=(10,5))
sns.lineplot(data=albert_prompt_epochs_lr5_bs8, x="epoch", y="train_accuracy", label="Train Accuracy - lr5_bs8")
sns.lineplot(data=albert_prompt_epochs_lr5_bs8, x="epoch", y="val_accuracy", label="Val Accuracy - lr5_bs8")
sns.lineplot(data=albert_prompt_epochs_lr5_bs16, x="epoch", y="train_accuracy", label="Train Accuracy - lr5_bs16")
sns.lineplot(data=albert_prompt_epochs_lr5_bs16, x="epoch", y="val_accuracy", label="Val Accuracy - lr5_bs16")
sns.lineplot(data=albert_prompt_epochs_lr1_bs8, x="epoch", y="train_accuracy", label="Train Accuracy - lr1_bs8")
sns.lineplot(data=albert_prompt_epochs_lr1_bs8, x="epoch", y="val_accuracy", label="Val Accuracy - lr1_bs8")
sns.lineplot(data=albert_prompt_epochs_lr1_bs16, x="epoch", y="train_accuracy", label="Train Accuracy - lr1_bs16")
sns.lineplot(data=albert_prompt_epochs_lr1_bs16, x="epoch", y="val_accuracy", label="Val Accuracy - lr1_bs16")
plt.title("Learning Curve: Accuracy - ALBERT w/ Prompt Tuning on Sentiment140")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.grid(True)
plt.show()

# All Prompt Tuning Training and Validation Loss
plt.figure(figsize=(10,5))
sns.lineplot(data=albert_prompt_epochs_lr5_bs8, x="epoch", y="train_loss", label="Train Loss - lr5_bs8")
sns.lineplot(data=albert_prompt_epochs_lr5_bs8, x="epoch", y="val_loss", label="Val Loss - lr5_bs8")
sns.lineplot(data=albert_prompt_epochs_lr5_bs16, x="epoch", y="train_loss", label="Train Loss - lr5_bs16")
sns.lineplot(data=albert_prompt_epochs_lr5_bs16, x="epoch", y="val_loss", label="Val Loss - lr5_bs16")
sns.lineplot(data=albert_prompt_epochs_lr1_bs8, x="epoch", y="train_loss", label="Train Loss - lr1_bs8")
sns.lineplot(data=albert_prompt_epochs_lr1_bs8, x="epoch", y="val_loss", label="Val Loss - lr1_bs8")
sns.lineplot(data=albert_prompt_epochs_lr1_bs16, x="epoch", y="train_loss", label="Train Loss - lr1_bs16")
sns.lineplot(data=albert_prompt_epochs_lr1_bs16, x="epoch", y="val_loss", label="Val Loss - lr1_bs16")
plt.title("Learning Curve: Loss - ALBERT w/ Prompt Tuning on Sentiment140")
```

```
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()
```



```
# Best Prompt Tuning Train/Val Acc Learning Curve

albert_prompt_epochs_map = {
    (5, 8): albert_prompt_epochs_lr5_bs8,
    (5, 16): albert_prompt_epochs_lr5_bs16,
    (1, 8): albert_prompt_epochs_lr1_bs8,
    (1, 16): albert_prompt_epochs_lr1_bs16
}

prompt_lr_mapping = {
    5e-5: "5e-5",
    1e-4: "1e-4"
}

prompt_best_lr_tag_for_map = {5e-5: 5, 1e-4: 1}[prompt_best_lr]

prompt_best_bs_tag = prompt_best_bs

prompt_epochs = albert_prompt_epochs_map[(prompt_best_lr_tag_for_map, prompt_best_bs_tag)]

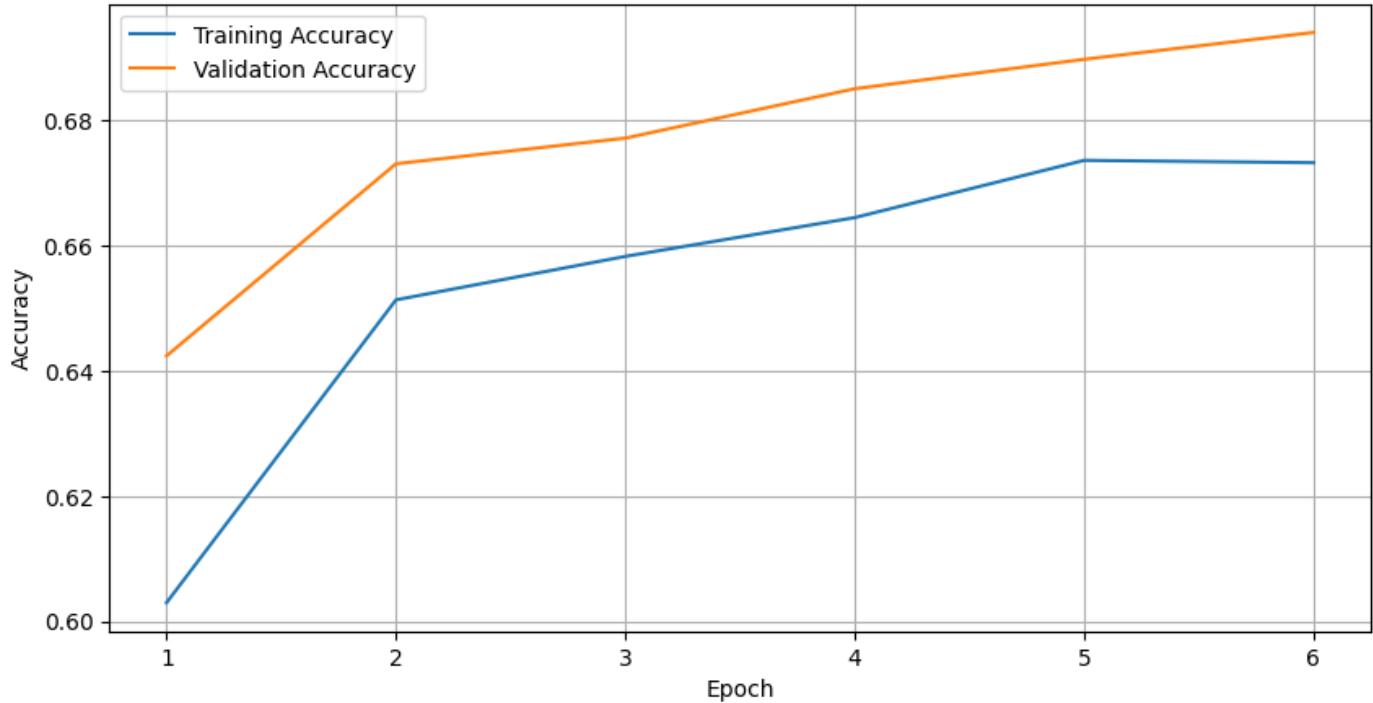
# Best Prompt Tuning Training and Validation Accuracy
plt.figure(figsize=(10,5))
sns.lineplot(data=prompt_epochs, x="epoch", y="train_accuracy", label="Training Accuracy")
sns.lineplot(data=prompt_epochs, x="epoch", y="val_accuracy", label="Validation Accuracy")
plt.title(f"Learning Curve: Best Accuracy - ALBERT w/ Prompt Tuning on Sentiment140")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.grid(True)
plt.show()

# Best Prompt Tuning Training and Validation Loss
plt.figure(figsize=(10,5))
```

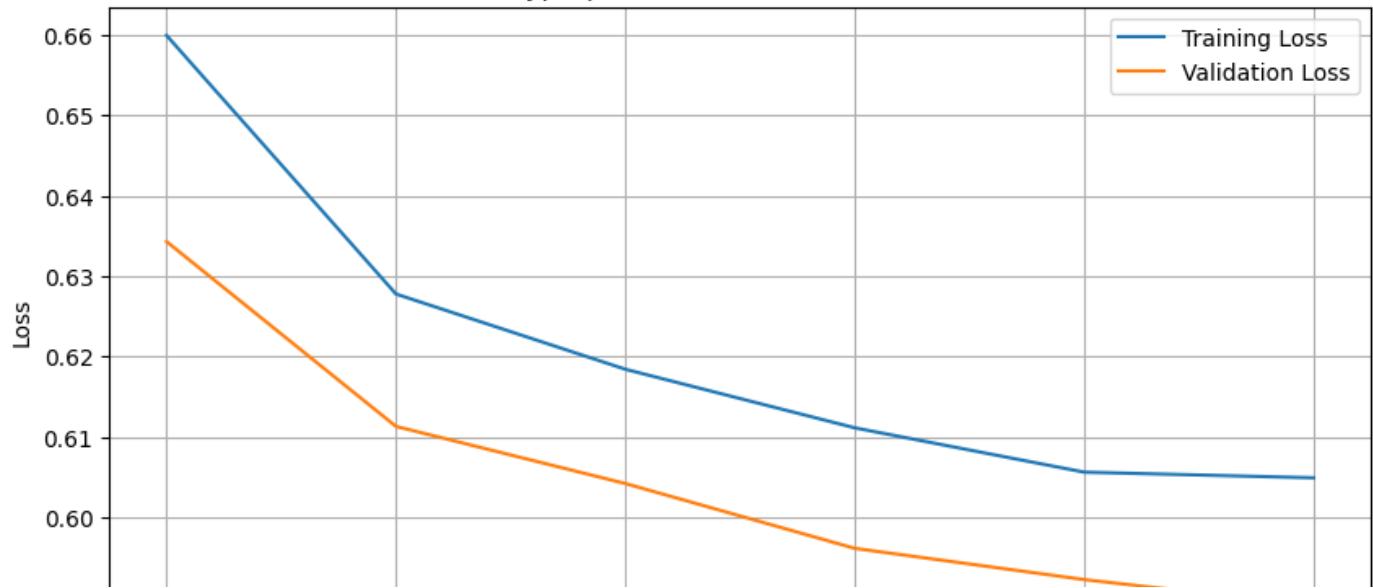
```
sns.lineplot(data=prompt_epochs, x="epoch", y="train_loss", label="Training Loss")
sns.lineplot(data=prompt_epochs, x="epoch", y="val_loss", label="Validation Loss")
plt.title(f"Learning Curve: Best Loss – ALBERT w/ Prompt Tuning on Sentiment140\n")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()
```



Learning Curve: Best Accuracy - ALBERT w/ Prompt Tuning on Sentiment140
Best Hyperparameters: LR: 0.0001 and BS: 8



Learning Curve: Best Loss - ALBERT w/ Prompt Tuning on Sentiment140
Best Hyperparameters: LR: 0.0001 and BS: 8





PEFT Method Comparison:

```
# Final results per BitFit/LoRA/Prompt Tuning Implementation

albert_bf_results = pd.read_csv('/content/sent_albert_bitfit_results.csv')
albert_lora_results = pd.read_csv('/content/sent_albert_lora_results.csv')
albert_prompt_results = pd.read_csv('/content/sent_albert_prompt_results.csv')

# Table of comparisons
comparison = pd.DataFrame({
    "Method": ["BitFit", "LoRA", "Prompt Tuning"],
    "Best Validation F1": [
        albert_bf_results["f1"].max(),
        albert_lora_results["f1"].max(),
        albert_prompt_results["f1"].max()
    ],
    "Best Validation Accuracy": [
        albert_bf_results["accuracy"].max(),
        albert_lora_results["accuracy"].max(),
        albert_prompt_results["accuracy"].max()
    ],
    "Runtime (sec)": [
        albert_bf_results["training_time"].sum(),
        albert_lora_results["training_time"].sum(),
        albert_prompt_results["training_time"].sum()
    ],
    "Inference Time (sec)": [

```

```

        albert_bf_results["inference_time"].sum(),
        albert_lora_results["inference_time"].sum(),
        albert_prompt_results["inference_time"].sum()
    ],
    "Max GPU Memory (GB)": [
        albert_bf_results["max_memory"].max(),
        albert_lora_results["max_memory"].max(),
        albert_prompt_results["max_memory"].max()
    ]
}

print("\nFinal Validation Performance PEFT Comparison - ALBERT on Sentiment140:")
display(comparison)

```



Final Validation Performance PEFT Comparison - ALBERT on Sentiment140:

	Method	Best Validation F1	Best Validation Accuracy	Runtime (sec)	Inference Time (sec)	Max GPU Memory (GB)
0	BitFit	0.799496	0.7995	3735.495723	57.115629	1.937313
1	LoRA	0.765120	0.7658	4097.389168	54.110241	1.937313

Next steps:

[Generate code with comparison](#)

[View recommended plots](#)

[New interactive sheet](#)

```

# Load overall results where inference_time is stored
albert_prompt_results = pd.read_csv('/content/sent_albert_prompt_results.csv')
albert_bf_results = pd.read_csv('/content/sent_albert_bitfit_results.csv')
albert_lora_results = pd.read_csv('/content/sent_albert_lora_results.csv')

# Manually map best learning rates to filename tags (from before)
lr_tag_mapping = {
    5e-5: "5e-05",
    1e-4: "0.0001"
}
bf_best_lr_tag = lr_tag_mapping[bf_best_lr]
lora_best_lr_tag = lr_tag_mapping[lora_best_lr]
prompt_best_lr_tag = lr_tag_mapping[prompt_best_lr]

# Load best inference metric summaries
bf_inf = pd.read_csv(f'/content/sent_albert_bitfit_inference_metrics_summary_lr{bf_lora_inf = pd.read_csv(f'/content/sent_albert_lora_inference_metrics_summary_lr{lor

```

```
prompt_inf = pd.read_csv(f'/content/sent_albert_prompt_inference_metrics_summary_lr

# Extract inference times
bf_inference_time = albert_bf_results[
    (albert_bf_results["learning_rate"] == bf_best_lr) &
    (albert_bf_results["batch_size"] == bf_best_bs)
]["inference_time"].values[0]

lora_inference_time = albert_lora_results[
    (albert_lora_results["learning_rate"] == lora_best_lr) &
    (albert_lora_results["batch_size"] == lora_best_bs)
]["inference_time"].values[0]

prompt_inference_time = albert_prompt_results[
    (albert_prompt_results["learning_rate"] == prompt_best_lr) &
    (albert_prompt_results["batch_size"] == prompt_best_bs)
]["inference_time"].values[0]

# Table of best per-implementation metrics (based on best lr and bs per PEFT method
final_test_results = pd.DataFrame({
    "Method": ["BitFit", "LoRA", "Prompt Tuning"],
    "Test Accuracy": [
        bf_inf.loc["accuracy", "precision"],
        lora_inf.loc["accuracy", "precision"],
        prompt_inf.loc["accuracy", "precision"]
    ],
    "F1 Macro": [
        bf_inf.loc["macro avg", "f1-score"],
        lora_inf.loc["macro avg", "f1-score"],
        prompt_inf.loc["macro avg", "f1-score"]
    ],
    "F1 Weighted": [
        bf_inf.loc["weighted avg", "f1-score"],
        lora_inf.loc["weighted avg", "f1-score"],
        prompt_inf.loc["weighted avg", "f1-score"]
    ],
    "Inference Time (sec)": [
        bf_inference_time,
        lora_inference_time,
        prompt_inference_time
    ]
})

print("\nFinal Test Set Inference Performance PEFT Comparison - ALBERT on Sentiment
display(final_test_results)
```



Final Test Set Inference Performance PEFT Comparison - ALBERT on Sentiment140:

Method	Test Accuracy	F1 Macro	F1 Weighted	Inference Time (sec)
0 BitFit	0.7995	0.799496	0.799496	9.836762
1 LoRA	0.7658	0.765120	0.765120	17.667457
2 Prompt	0.6941	0.693660	0.693660	640.500990

Next steps:

[Generate code with final_test_results](#)[View recommended plots](#)[New interactive](#)

```
# Zip the entire /content folder
!zip -r /content/ALBERT_Sentiment_solo.zip /content
```

```
# Download the zipped file
from google.colab import files
files.download('/content/ALBERT_Sentiment_solo.zip')
```

→ adding: content/ (stored 0%)
 adding: content/.config/ (stored 0%)
 adding: content/.config/.last_opt_in_prompt.yaml (stored 0%)
 adding: content/.config/active_config (stored 0%)
 adding: content/.config/config_sentinel (stored 0%)
 adding: content/.config/.last_survey_prompt.yaml (stored 0%)
 adding: content/.config/configurations/ (stored 0%)
 adding: content/.config/configurations/config_default (deflated 15%)
 adding: content/.config/hidden_gcloud_config_universe_descriptor_data_cache_ (deflated 15%)
 adding: content/.config/logs/ (stored 0%)
 adding: content/.config/logs/2025.04.24/ (stored 0%)
 adding: content/.config/logs/2025.04.24/18.19.38.522066.log (deflated 58%)
 adding: content/.config/logs/2025.04.24/18.19.46.929623.log (deflated 87%)
 adding: content/.config/logs/2025.04.24/18.19.56.709493.log (deflated 57%)
 adding: content/.config/logs/2025.04.24/18.19.57.353004.log (deflated 56%)
 adding: content/.config/logs/2025.04.24/18.19.48.089267.log (deflated 58%)
 adding: content/.config/logs/2025.04.24/18.19.17.922226.log (deflated 93%)
 adding: content/.config/default_configs.db (deflated 98%)
 adding: content/.config/gce (stored 0%)
 adding: content/.config/.last_update_check.json (deflated 23%)
 adding: content/sent_albert_lora_epoch_logs_lr5e-05_bs8.csv (deflated 47%)
 adding: content/sent_albert_prompt_inference_metrics_summary_lr0.0001_bs16.csv (deflated 47%)
 adding: content/sent_albert_lora_inference_metrics_summary_lr5e-05_bs8.csv (deflated 52%)
 adding: content/sent_albert_bitfit_inference_metrics_summary_lr5e-05_bs16.csv (deflated 52%)
 adding: content/sent_albert_lora_epoch_logs_lr5e-05_bs16.csv (deflated 52%)
 adding: content/sent_albert_lora_epoch_logs_lr0.0001_bs16.csv (deflated 52%)

```
adding: content/sent_albert_bitfit_epoch_logs_lr5e-05_bs16.csv (deflated 41%)
adding: content/sent_albert_prompt_inference_predictions_lr5e-05_bs16.csv (deflated 41%)
adding: content/sent_albert_bitfit_epoch_logs_lr0.0001_bs8.csv (deflated 47%)
adding: content/sent_albert_prompt_inference_metrics_summary_lr5e-05_bs16.csv (deflated 47%)
adding: content/sent_albert_bitfit_inference_predictions_lr5e-05_bs16.csv (deflated 47%)
adding: content/sent_albert_bitfit_inference_predictions_lr5e-05_bs8.csv (deflated 47%)
adding: content/sentiment140.py (deflated 62%)
adding: content/sent_albert_bitfit_inference_metrics_summary_lr0.0001_bs8.csv (deflated 47%)
adding: content/sent_albert_lora_inference_predictions_lr5e-05_bs8.csv (deflated 47%)
adding: content/sent_albert_prompt_epoch_logs_lr0.0001_bs16.csv (deflated 49%)
adding: content/sent_albert_prompt_final_comparison_prompt_tuning.csv (deflated 49%)
adding: content/sent_albert_lora_final_comparison_lora.csv (deflated 18%)
adding: content/sent_albert_lora_inference_predictions_lr0.0001_bs16.csv (deflated 49%)
adding: content/sent_albert_prompt_epoch_logs_lr5e-05_bs8.csv (deflated 48%)
adding: content/sent_albert_bitfit_inference_metrics_summary_lr0.0001_bs16.csv (deflated 48%)
adding: content/sent_albert_prompt_inference_predictions_lr5e-05_bs8.csv (deflated 48%)
adding: content/sent_albert_lora_inference_metrics_summary_lr5e-05_bs16.csv (deflated 48%)
adding: content/sent_albert_bitfit_inference_predictions_lr0.0001_bs16.csv (deflated 48%)
adding: content/sent_albert_prompt_inference_predictions_lr0.0001_bs16.csv (deflated 48%)
adding: content/sent_albert_bitfit_results.csv (deflated 47%)
adding: content/sent_albert_lora_inference_predictions_lr5e-05_bs16.csv (deflated 47%)
adding: content/sent_albert_prompt_results.csv (deflated 49%)
adding: content/sent_albert_bitfit_inference_metrics_summary_lr5e-05_bs8.csv (deflated 48%)
adding: content/sent_albert_prompt_inference_metrics_summary_lr5e-05_bs8.csv (deflated 48%)
adding: content/sent_albert_bitfit_epoch_logs_lr0.0001_bs16.csv (deflated 48%)
adding: content/sent_albert_prompt_inference_predictions_lr0.0001_bs8.csv (deflated 48%)
adding: content/sent_albert_lora_epoch_logs_lr0.0001_bs8.csv (deflated 52%)
adding: content/sent_albert_lora_inference_metrics_summary_lr0.0001_bs16.csv (deflated 52%)
adding: content/sent_albert_bf_final_comparison_bitfit.csv (deflated 19%)
adding: content/sent_albert_bitfit_inference_predictions_lr0.0001_bs8.csv (deflated 46%)
adding: content/sent_albert_lora_results.csv (deflated 46%)
adding: content/sent_albert_lora_inference_predictions_lr0.0001_bs8.csv (deflated 46%)
adding: content/sent_albert_prompt_inference_metrics_summary_lr0.0001_bs8.csv (deflated 46%)
adding: content/sent_albert_lora_inference_metrics_summary_lr0.0001_bs8.csv (deflated 46%)
adding: content/sent_albert_prompt_epoch_logs_lr5e-05_bs16.csv (deflated 49%)
adding: content/sent_albert_prompt_epoch_logs_lr0.0001_bs8.csv (deflated 49%)
adding: content/sent_albert_bitfit_epoch_logs_lr5e-05_bs8.csv (deflated 48%)
adding: content/sample_data/ (stored 0%)
adding: content/sample_data/anscombe.json (deflated 83%)
adding: content/sample_data/README.md (deflated 39%)
adding: content/sample_data/mnist_test.csv (deflated 88%)
adding: content/sample_data/mnist_train_small.csv (deflated 88%)
adding: content/sample_data/california_housing_train.csv (deflated 79%)
adding: content/sample_data/california_housing_test.csv (deflated 76%)
```

